

Assignment Project Exam Help

5QQMN534: Algorithmic Finance

<https://tutorcs.com>
WeChat: cstutorcs

Week3: Mastering Basics – Financial Data Analysis with
Pandas

Yves Hilpisch - Python for Finance 2nd Edition 2019: Chapter 5

Agenda

- Data Analysis with Pandas
- The DataFrame Class
 - First Steps with a DataFrame Class
 - * .loc and .iloc functions
 - Second Steps with a DataFrame Class
- Basic Analytics
- Basic Visualisation
- The Series Class
- GroupBy Operations
- Complex Selection
- Concatenation, Joining and Merging
- Performance Aspects
- Conclusion

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Data Analysis with pandas1

- This chapter is about pandas, a library for data analysis with a focus on tabular data. pandas is a powerful tool that not only provides many useful classes and functions but also does a great job of wrapping functionality from other packages.
- The result is a user interface that makes data analysis, and in particular financial analysis, a convenient and efficient task. This chapter covers the following fundamental data structures:

| Object type | Meaning | Used for |
|-------------|--------------------------------------|-----------------------------------|
| DataFrame | 2-dimensional data object with index | Tabular data organized in columns |
| Series | 1-dimensional data object with index | Single (time) series of data |

The chapter is organized as follows: WeChat: cstutorcs

“The DataFrame Class”

This section starts by exploring the basic characteristics and capabilities of the DataFrame class of pandas by using simple and small data sets; it then shows how to transform a NumPy ndarray object into a DataFrame object.

“Basic Analytics” and “Basic Visualization”

Basic analytics and visualization capabilities are introduced in these sections (later chapters go deeper into these topics).

Data Analysis with pandas2

“The Series Class”

This rather brief section covers the `Series` class of `pandas`, which in a sense represents a special case of the `DataFrame` class with a single column of data only.

“GroupBy Operations”

One of the strengths of the `DataFrame` class lies in grouping data according to a single or multiple columns. This section explores the grouping capabilities of `pandas`.

<https://tutorcs.com>

“Complex Selection”

This section illustrates how the use of (complex) conditions allows for the easy selection of data from a `DataFrame` object.

WeChat: estutorcs

“Concatenation, Joining, and Merging”

The combining of different data sets into one is an important operation in data analysis. `pandas` provides different options to accomplish this task, as described in this section.

“Performance Aspects”

Like Python in general, `pandas` often provides multiple options to accomplish the same goal. This section takes a brief look at potential performance differences.

The DataFrame Class

- At the core of pandas (and this chapter) is the DataFrame, a class designed to efficiently handle data in tabular form — i.e., data characterized by a columnar organization.
- To this end, the DataFrame class provides, for instance, column labelling as well as flexible indexing capabilities for the rows (records) of the data set, similar to a table in a relational database or an Excel spreadsheet.
- This section covers some fundamental aspects of the pandas DataFrame class.
- The class is so complex and powerful that only a fraction of its capabilities can be presented here.
- Subsequent chapters provide more examples and shed light on different aspects.

Documentation: <https://pandas.pydata.org/pandas-docs/stable/reference/frame.html>

First Steps with a DataFrame Class1

- On a fundamental level, the `DataFrame` class is designed to manage indexed and labelled data, not too different from a SQL database table or a worksheet in a spreadsheet application.
- Consider the following creation of a `DataFrame` object ❶

```
In [1]: import pandas as pd ❶
```

Imports pandas.

```
In [2]: df = pd.DataFrame([10, 20, 30, 40], ❷  
    columns=['numbers'], ❸  
    index=['a', 'b', 'c', 'd']) ❹
```

Defines the data as a list object.

Specifies the column label.

Specifies the index values/labels.

```
In [3]: df ❺
```

```
Out[3]:
```

| | numbers |
|---|---------|
| a | 10 |
| b | 20 |
| c | 30 |
| d | 40 |

Shows the data as well as column and index labels of the `DataFrame` object.

- This simple example already shows some major features of the `DataFrame` class when it comes to storing data:
- Data itself can be provided in different shapes and types (`list`, `tuple`, `ndarray`, and `dict` objects are candidates).
- Data is organized in columns, which can have custom names (labels).
- There is an index that can take on different formats (e.g., numbers, strings, time information).

First Steps with a DataFrame Class2

- Working with a DataFrame object is in general pretty convenient and efficient with regard to the handling of the object, e.g., compared to regular ndarray objects, which are more specialized and more restricted when one wants to (say) enlarge an existing object.
- At the same time, DataFrame objects are often as computationally efficient as ndarray objects.
- The following are simple examples showing how typical operations on a DataFrame object work:

❶

The `index` attribute and `Index` object.

❷

The `columns` attribute and `Index` object.

❸

Selects the value corresponding to index `c`.

❹

Selects the two values corresponding to indices `a` and `d`.

❺

Selects the second and third rows via the index positions.

❻

Calculates the sum of the single column.

❼

Uses the `apply()` method to calculate squares in vectorized fashion.

❽

Applies vectorization directly as with ndarray objects.

```
In [4]: df.index ❶  
Out[4]: Index(['a', 'b', 'c', 'd'], dtype='object')
```

```
In [5]: df.columns ❷  
Out[5]: Index(['numbers'], dtype='object')
```

```
In [6]: df.loc['c'] ❸  
Out[6]: numbers      30  
        Name: c, dtype: int64
```

```
In [7]: df.loc[['a', 'd']] ❹  
Out[7]:      numbers  
a         10  
d         40
```

```
In [8]: df.iloc[1:3] ❺  
Out[8]:      numbers  
b         20  
c         30
```

```
In [9]: df.sum() ❻  
Out[9]: numbers      100  
        dtype: int64
```

```
In [10]: df.apply(lambda x: x ** 2) ❼  
Out[10]:      numbers  
a         100  
b         400  
c         900  
d        1600
```

```
In [11]: df ** 2 ❽  
Out[11]:      numbers  
a         100  
b         400  
c         900  
d        1600
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

*.loc and .iloc functions

- loc gets rows (or columns) with **particular** labels from the index. Can use DateTime Indexes
- iloc gets rows (or columns) at particular positions in the index (so it **only** takes integers).

Python Pandas Selections and Indexing

Assignment Project Exam Help

.iloc selections - position based selection

`data.iloc[<row selection>, <column selection>]`

Integer list of rows: [0,1,2]
Slice of rows: [4:7]
Single values: 1

Integer list of columns: [0,1,2]
Slice of columns: [4:7]
Single column selections: 1

loc selections - position based selection

`data.loc[<row selection>, <column selection>]`

Index/Label value: 'john'
List of labels: ['john', 'sarah']
Logical/Boolean index: data['age'] == 10

Named column: 'first_name'
List of column names: ['first_name', 'age']
Slice of columns: 'first_name':'address'

First Steps with a DataFrame Class3

- Contrary to NumPy ndarray objects, enlarging the DataFrame object in both dimensions is possible:

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

```
In [12]: df['floats'] = (1.5, 2.5, 3.5, 4.5) ❶
```

```
In [13]: df
```

```
Out[13]:
```

| | numbers | floats |
|---|---------|--------|
| a | 10 | 1.5 |
| b | 20 | 2.5 |
| c | 30 | 3.5 |
| d | 40 | 4.5 |

❶

Adds a new column with float objects provided as a tuple object.

❷

Selects this column and shows its data and index labels.

```
In [14]: df['floats'] ❷
```

```
Out[14]:
```

| | |
|---|-----|
| a | 1.5 |
| b | 2.5 |
| c | 3.5 |
| d | 4.5 |

Name: floats, dtype: float64

First Steps with a DataFrame Class4

- A whole DataFrame object can also be taken to define a new column.
- In such a case, indices are aligned automatically:

Assignment Project Exam Help

```
In [15]: df['names'] = pd.DataFrame(['Yves', 'Sandra', 'Lilli', 'Henry'],  
                                   index=['a', 'b', 'c', 'd'])
```

```
In [16]: df
```

```
Out[16]:
```

| | numbers | floats | names |
|---|---------|--------|--------|
| a | 10 | 1.5 | Sandra |
| b | 20 | 2.5 | Lilli |
| c | 30 | 3.5 | Henry |
| d | 40 | 4.5 | Yves |

WeChat: cstutorcs

①

Another new column is created based on a DataFrame object.

First Steps with a DataFrame Class

- Appending data works similarly.
- However, in the following example a side effect is seen that is usually to be avoided - namely, the index gets replaced by a simple range index:

```
In [17]: df.append({'numbers': 100, 'floats': 5.75, 'names': 'Jil'},
                  ignore_index=True) ❶
```

```
Out[17]:
```

| | numbers | floats | names |
|---|---------|--------|--------|
| 0 | 10 | 1.50 | Sandra |
| 1 | 20 | 2.50 | Lilli |
| 2 | 30 | 3.50 | Henry |
| 3 | 40 | 4.50 | Yves |
| 4 | 100 | 5.75 | Jil |

```
In [18]: df = df.append(pd.DataFrame({'numbers': 100, 'floats': 5.75,
                                     'names': 'Jil'}, index=['y',])) ❷
```

```
In [19]: df
Out[19]:
```

| | numbers | floats | names |
|---|---------|--------|--------|
| a | 10 | 1.50 | Sandra |
| b | 20 | 2.50 | Lilli |
| c | 30 | 3.50 | Henry |
| d | 40 | 4.50 | Yves |
| | 100 | 5.75 | Jil |

```
In [20]: df = df.append(pd.DataFrame({'names': 'Liz'}, index=['z',]),
                       sort=False) ❸
```

```
In [21]: df
Out[21]:
```

| | numbers | floats | names |
|---|---------|--------|--------|
| a | 10.0 | 1.50 | Sandra |
| b | 20.0 | 2.50 | Lilli |
| c | 30.0 | 3.50 | Henry |
| d | 40.0 | 4.50 | Yves |
| y | 100.0 | 5.75 | Jil |

❶

Appends a new row via a dict object; this is a temporary operation during which index information gets lost.

❷

Appends the row based on a DataFrame object with index information; the original index information is preserved.

❸

Appends an incomplete data row to the DataFrame object, resulting in NaN values.

❹

Returns the different dtypes of the single columns; this is similar to what's possible with structured ndarray objects.

```
dtype: object
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

First Steps with a DataFrame Class6

- Although there are now missing values, the majority of method calls will still work:

```
In [23]: df[['numbers', 'floats']].mean()
Out[23]: numbers    10.00
         floats      3.55
         dtype: float64
```

```
In [24]: df[['numbers', 'floats']].std()
Out[24]: numbers    35.355339
         floats     1.662077
         dtype: float64
```

①

Calculates the mean over the two columns specified (ignoring rows with NaN values).

②

Calculates the standard deviation over the two columns specified (ignoring rows with NaN values).

Second Steps with a DataFrame Class1

- The example in this subsection is based on an `ndarray` object with standard normally distributed random numbers.
- It explores further features such as a `DatetimeIndex` to manage time series data:

Assignment Project Exam Help

```
In [25]: import numpy as np
```

```
In [26]: np.random.seed(1)
```

```
In [27]: a = np.random.standard_normal((9, 4))
```

```
In [28]: a
```

```
Out[28]: array([[ -1.74976547,  0.3426804 ,  1.1530358 , -0.25243604],
 [  0.98132079,  0.51421884,  0.22117967, -1.07004333],
 [-0.18949583,  0.25500144, -0.45802699,  0.43516349],
 [-0.58359505,  0.81684707,  0.67272081, -0.10441114],
 [-0.53128038,  1.02973269, -0.43813562, -1.11831825],
 [  1.61898166,  1.54160517, -0.25187914, -0.84243574],
 [  0.18451869,  0.9370822 ,  0.73100034,  1.36155613],
 [-0.32623806,  0.05567601,  0.22239961, -1.443217  ],
 [-0.75635231,  0.81645401,  0.75044476, -0.45594693]])
```

<https://tutorcs.com>
WeChat: cstutorcs

Second Steps with a DataFrame Class2

- Although one can construct `DataFrame` objects more directly (as seen before), using an `ndarray` object is generally a good choice since `pandas` will retain the basic structure and will “only” add meta-information (e.g., index values).
- It also represents a typical use case for financial applications and scientific research in general.
- For example:

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

```
In [29]: df = pd.DataFrame(a) ❶
```

```
In [30]: df
```

```
Out[30]:
```

| | 0 | 1 | 2 | 3 |
|---|-----------|----------|-----------|-----------|
| 0 | -1.749765 | 0.342680 | 1.153036 | -0.252436 |
| 1 | 0.981321 | 0.514219 | 0.221180 | -1.070043 |
| 2 | -0.189496 | 0.255001 | -0.458027 | 0.435163 |
| 3 | -0.583595 | 0.816847 | 0.672721 | -0.104411 |
| 4 | -0.531280 | 1.029733 | -0.438136 | -1.118318 |
| 5 | 1.618982 | 1.541605 | -0.251879 | -0.842436 |
| 6 | 0.184519 | 0.937082 | 0.731000 | 1.361556 |
| 7 | -0.326238 | 0.055676 | 0.222400 | -1.443217 |
| 8 | -0.756352 | 0.816454 | 0.750445 | -0.455947 |

Creates a `DataFrame` object from the `ndarray` object.

Second Steps with a DataFrame Class3

- **Table 5-1** lists the parameters that the `DataFrame()` function takes. In the table, “array-like” means a data structure similar to an `ndarray` object — a `list`, for example. `Index` is an instance of the `pandas Index` class.

Assignment Project Exam Help

Table 5-1. Parameters of `DataFrame()` function

| Parameter | Format | Description |
|----------------------|---|--|
| <code>data</code> | <code>ndarray</code> / <code>dict</code> / <code>DataFrame</code> | Data for <code>DataFrame</code> , <code>dict</code> can contain <code>Series</code> , <code>ndarray</code> , <code>list</code> |
| <code>index</code> | <code>Index</code> / <code>array-like</code> | Index to use; defaults to <code>range(n)</code> |
| <code>columns</code> | <code>Index</code> / <code>array-like</code> | Column headers to use; defaults to <code>range(n)</code> |
| <code>dtype</code> | <code>dtype</code> , default <code>None</code> | Data type to use/force; otherwise, it is inferred |
| <code>copy</code> | <code>bool</code> , default <code>None</code> | Copy data from inputs |

Second Steps with a DataFrame Class4

- As with structured arrays, and as seen before, `DataFrame` objects have column names that can be defined directly by assigning a `list` object with the right number of elements.
- This illustrates that one can define/change the attributes of the `DataFrame` object easily:

```
In [31]: df.columns = ['No1', 'No2', 'No3', 'No4']
```

```
In [32]: df
```

```
Out[32]:
```

| | No1 | No2 | No3 | No4 |
|---|-----------|----------|-----------|-----------|
| 0 | -1.749765 | 0.342680 | 1.133063 | -1.352486 |
| 1 | 0.981321 | 0.514219 | 0.221180 | -1.070043 |
| 2 | -0.189496 | 0.255001 | -0.458027 | 0.435163 |
| 3 | -0.583595 | 0.816847 | 0.672721 | -0.104411 |
| 4 | -0.531280 | 1.029733 | -0.438136 | -1.118318 |
| 5 | 1.618982 | 1.541605 | -0.251879 | -0.842436 |
| 6 | 0.184519 | 0.937082 | 0.731000 | 1.361556 |
| 7 | -0.326238 | 0.055676 | 0.222400 | -1.443217 |
| 8 | -0.756352 | 0.816454 | 0.750445 | -0.455947 |

```
In [33]: df['No2'].mean()
Out[33]: 0.7010330941456459
```

①

Specifies the column labels via a `list` object.

②

Picking a column is now made easy.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estutorcs

Second Steps with a DataFrame Class5

- To work with financial time series data efficiently, one must be able to handle time indices well.
- This can also be considered a major strength of pandas.
- For example, assume that our nine data entries in the four columns correspond to month-end data, beginning in January 2019.
- A `DatetimeIndex` object is then generated with the `date_range()` function as follows:

```
In [34]: dates = pd.date_range('2019-1-1', periods=9, freq='M') ❶

In [35]: dates
Out[35]: DatetimeIndex(['2019-01-31', '2019-02-28', '2019-03-31', '2019-04-30',
                        '2019-05-31', '2019-06-30', '2019-07-31', '2019-08-31',
                        '2019-09-30'],
                        dtype='datetime64[ns]', freq='M')
```

❶

Creates a `DatetimeIndex` object.

Second Steps with a DataFrame Class6

Table 5-2 lists the parameters that the `date_range()` function takes.

Table 5-2. Parameters of `date_range()` function

| Parameter | Format | Description |
|------------------------|----------------------|---|
| <code>start</code> | string/datetime | Left bound for generating dates |
| <code>end</code> | string/datetime | Right bound for generating dates |
| <code>periods</code> | integer/None | Number of periods (if <code>start</code> or <code>end</code> is None) |
| <code>freq</code> | string/DateOffset | Frequency string, e.g., 5D for 5 days |
| <code>tz</code> | string/None | Time zone name for localized index |
| <code>normalize</code> | bool, default None | Normalizes <code>start</code> and <code>end</code> to midnight |
| <code>name</code> | string, default None | Name of resulting index |

Second Steps with a DataFrame Class7

- The following code defines the just-created `DatetimeIndex` object as the relevant index object, making a time series of the original data set:

Assignment Project Exam Help

`In [36]: df.index = dates`

`In [37]: df`

`Out[37]:`

| | No1 | No2 | No3 | No4 |
|------------|-----------|----------|-----------|-----------|
| 2019-01-31 | -1.749785 | 0.342680 | 1.153036 | -0.252436 |
| 2019-02-28 | 0.981321 | 0.514219 | 0.221180 | -1.070043 |
| 2019-03-31 | -0.189496 | 0.255001 | -0.458027 | 0.435163 |
| 2019-04-30 | -0.583595 | 0.816847 | 0.672721 | -0.104411 |
| 2019-05-31 | -0.531280 | 1.029733 | -0.438136 | -1.118318 |
| 2019-06-30 | 1.618982 | 1.541605 | -0.251879 | -0.842436 |
| 2019-07-31 | 0.184519 | 0.937082 | 0.731000 | 1.361556 |
| 2019-08-31 | -0.326238 | 0.055676 | 0.222400 | -1.443217 |
| 2019-09-30 | -0.756352 | 0.816454 | 0.750445 | -0.455947 |

<https://tutorcs.com>

WeChat: cstutorcs

Second Steps with a DataFrame Class8

Table 5-3. Frequency parameter values
for `date_range()` function

| Alias | Description |
|-------|--|
| B | Business day frequency |
| C | Custom business day frequency (experimental) |
| D | Calendar day frequency |
| W | Weekly frequency |
| M | Month end frequency |
| BM | Business month end frequency |
| MS | Month start frequency |
| BMS | Business month start frequency |
| Q | Quarter end frequency |
| BQ | Business quarter end frequency |

| Alias | Description |
|-------|----------------------------------|
| QS | Quarter start frequency |
| BQS | Business quarter start frequency |
| A | Year end frequency |
| BA | Business year end frequency |
| AS | Year start frequency |
| BAS | Business year start frequency |
| H | Hourly frequency |
| T | Minutely frequency |
| S | Secondly frequency |
| L | Milliseconds |
| U | Microseconds |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estutores

Second Steps with a DataFrame Class9

- In some circumstances, it pays off to have access to the original data set in the form of the `ndarray` object.
- The `values` attribute provides direct access to it:

```
In [38]: df.values
Out[38]: array([[ -1.74976547,  0.3426804 ,  1.1530358 , -0.25243604],
 [  0.98132079,  0.51421884,  0.22117967, -1.07004333],
 [-0.18949583,  0.25500144, -0.45802699,  0.43516349],
 [-0.58359505,  0.81684707,  0.67272081, -0.10441114],
 [-0.53128038,  1.02973269, -0.43813562, -1.11831825],
 [  1.61898166,  1.54160517, -0.25187914, -0.84243574],
 [  0.18451869,  0.9370822 ,  0.73100034,  1.36155613],
 [-0.32623806,  0.05567601,  0.22239961, -1.443217  ],
 [-0.75635231,  0.81645401,  0.75044476, -0.45594693]])
```

```
In [39]: np.array(df)
Out[39]: array([[ -1.74976547,  0.3426804 ,  1.1530358 , -0.25243604],
 [  0.98132079,  0.51421884,  0.22117967, -1.07004333],
 [-0.18949583,  0.25500144, -0.45802699,  0.43516349],
 [-0.58359505,  0.81684707,  0.67272081, -0.10441114],
 [-0.53128038,  1.02973269, -0.43813562, -1.11831825],
 [  1.61898166,  1.54160517, -0.25187914, -0.84243574],
 [  0.18451869,  0.9370822 ,  0.73100034,  1.36155613],
 [-0.32623806,  0.05567601,  0.22239961, -1.443217  ],
 [-0.75635231,  0.81645401,  0.75044476, -0.45594693]])
```

ARRAYS AND DATAFRAMES

One can generate a `DataFrame` object from an `ndarray` object, but one can also easily generate an `ndarray` object out of a `DataFrame` by using the `values` attribute of the `DataFrame` class or the function `np.array()` of NumPy.

Basic Analytics1

- Like the NumPy ndarray class, the pandas DataFrame class has a multitude of convenience methods built in.
- As a starter, consider the methods `info()` and `describe()`:

```
In [40]: df.info() ①
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 9 entries, 2019-01-31 to 2019-09-30
Freq: M
Data columns (total 4 columns):
No1      9 non-null float64
No2      9 non-null float64
No3      9 non-null float64
No4      9 non-null float64
dtypes: float64(4)
memory usage: 360.0 bytes
```

```
In [41]: df.describe() ②
Out[41]:
```

| | No1 | No2 | No3 | No4 |
|-------|-----------|----------|-----------|-----------|
| count | 9.000000 | 9.000000 | 9.000000 | 9.000000 |
| mean | -0.150212 | 0.701033 | 0.289193 | -0.387788 |
| std | 0.988306 | 0.457685 | 0.579920 | 0.877532 |
| min | -1.749765 | 0.055676 | -0.458027 | -1.443217 |
| 25% | -0.583595 | 0.342680 | -0.251879 | -1.070043 |
| 50% | -0.326238 | 0.816454 | 0.222400 | -0.455947 |
| 75% | 0.184519 | 0.937082 | 0.731000 | -0.104411 |
| max | 1.618982 | 1.541605 | 1.153036 | 1.361556 |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

This method prints information about a DataFrame including the index dtype and columns, non-null values and memory usage. Note the number of cells in each column (non-null values).

① Provides meta-information regarding the data, columns, and index.

② Provides helpful summary statistics per column (for numerical data).

25th Percentile - Also known as the first, or lower, quartile. The 25th percentile is the value at which 25% of the answers lie below that value, and 75% of the answers lie above that value.

50th Percentile - Also known as the Median. The median cuts the data set in half. Half of the answers lie below the median and half lie above the median.

75th Percentile - Also known as the third, or upper, quartile. The 75th percentile is the value at which 25% of the answers lie above that value and 75% of the answers lie below that value.

Basic Analytics2

- In addition, one can easily get the column-wise or row-wise sums, means, and cumulative sums:

```
In [43]: df.sum()
Out[43]: No1    -1.351906
         No2     6.309298
         No3     2.602739
         No4    -3.490089
         dtype: float64
```

```
In [44]: df.mean()
Out[44]: No1    -0.150212
         No2     0.701033
         No3     0.289193
         No4    -0.387788
         dtype: float64
```

```
In [45]: df.mean(axis=0)
Out[45]: No1    -0.150212
         No2     0.701033
         No3     0.289193
         No4    -0.387788
         dtype: float64
```

Column-wise sum.

```
In [46]: df.mean(axis=1)
Out[46]: 2019-01-31    -0.126621
         2019-02-28    -0.014639
         2019-03-31     0.010661
         2019-04-30     0.200390
         2019-05-31    -0.264500
         2019-06-30     0.019100
         2019-07-31     0.803539
         2019-08-31    -0.372845
         2019-09-30     0.088650
         Freq: M, dtype: float64
```

Column-wise mean.

Row-wise mean.

```
In [47]: df.cumsum()
Out[47]:
```

| | No1 | No2 | No3 | No4 |
|------------|-----------|----------|----------|-----------|
| 2019-01-31 | -1.749765 | 0.342680 | 1.153036 | -0.252436 |
| 2019-02-28 | -0.768445 | 0.856899 | 1.374215 | -1.322479 |
| 2019-03-31 | -0.957941 | 1.111901 | 0.916188 | -0.887316 |
| 2019-04-30 | -1.541536 | 1.928748 | 1.588909 | -0.991727 |
| 2019-05-31 | -2.072816 | 2.958480 | 1.150774 | -2.110045 |
| 2019-06-30 | -0.453834 | 4.500086 | 0.898895 | -2.952481 |
| 2019-07-31 | -0.269316 | 5.437168 | 1.629895 | -1.590925 |
| 2019-08-31 | -0.595554 | 5.492844 | 1.852294 | -3.034142 |
| 2019-09-30 | -1.351906 | 6.309298 | 2.602739 | -3.490089 |

Column-wise cumulative sum (starting at first index position).

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Basic Analytics3

- DataFrame objects also understand NumPy universal functions, as expected:

```
In [48]: np.mean(df)
Out[48]: No1    -0.150212
         No2     0.701033
         No3     0.289193
         No4    -0.387788
         dtype: float64
```

```
In [49]: np.log(df)
Out[49]:
```

| | No1 | No2 | No3 | No4 |
|------------|-----------|-----------|-----------|-----------|
| 2019-01-31 | NaN | -1.070957 | 0.142398 | NaN |
| 2019-02-28 | -0.019856 | -0.665106 | -1.508780 | NaN |
| 2019-03-31 | NaN | -1.366486 | NaN | -0.832033 |
| 2019-04-30 | NaN | -0.202303 | -0.396425 | NaN |
| 2019-05-31 | NaN | 0.029299 | NaN | NaN |
| 2019-06-30 | 0.481797 | 0.432824 | NaN | NaN |
| 2019-07-31 | -1.690005 | -0.064984 | -0.313341 | 0.308628 |
| 2019-08-31 | NaN | -2.888206 | -1.503279 | NaN |
| 2019-09-30 | NaN | -0.202785 | -0.287089 | NaN |

- Column-wise mean.
- Element-wise natural logarithm; a warning is raised but the calculation runs through, resulting in multiple NaN values.
- Element-wise square root for the absolute values ...

```
np.sqrt(abs(df))
```

| | No1 | No2 | No3 | No4 |
|------------|----------|----------|----------|----------|
| 2019-01-31 | 1.322787 | 0.585389 | 1.073795 | 0.502430 |
| 2019-02-28 | 0.990616 | 0.717091 | 0.470297 | 1.034429 |
| 2019-03-31 | 0.435311 | 0.504977 | 0.676777 | 0.659669 |
| 2019-04-30 | 0.763934 | 0.903796 | 0.820196 | 0.323127 |
| 2019-05-31 | 0.728890 | 1.014757 | 0.661918 | 1.057506 |
| 2019-06-30 | 1.272392 | 1.241614 | 0.501876 | 0.917843 |
| 2019-07-31 | 0.429556 | 0.968030 | 0.854986 | 1.166857 |
| 2019-08-31 | 0.571173 | 0.235958 | 0.471593 | 1.201340 |
| 2019-09-30 | 0.869685 | 0.903578 | 0.866282 | 0.675238 |

Basic Analytics4

```
In [51]: np.sqrt(abs(df)).sum()
```

```
Out[51]: No1      7.384345  
        No2      7.075190  
        No3      6.397719  
        No4      7.538440  
        dtype: float64
```

Assignment Project Exam Help

<https://tutorcs.com>

A linear transform of the numerical data.

```
In [52]: 100 * df + 100
```

```
Out[52]:
```

| | No1 | No2 | No3 | No4 |
|------------|------------|------------|------------|------------|
| 2019-01-31 | -74.976547 | 134.268040 | 215.303580 | 74.756396 |
| 2019-02-28 | 198.132079 | 151.421884 | 122.117967 | -7.004333 |
| 2019-03-31 | 81.050417 | 125.500144 | 54.197301 | 143.516349 |
| 2019-04-30 | 41.640495 | 181.684707 | 167.272081 | 89.558886 |
| 2019-05-31 | 46.871962 | 202.973269 | 56.186438 | -11.831825 |
| 2019-06-30 | 261.898166 | 254.160517 | 74.812086 | 15.756426 |
| 2019-07-31 | 118.451869 | 193.708220 | 173.100034 | 236.155613 |
| 2019-08-31 | 67.376194 | 105.567601 | 122.239961 | -44.321700 |
| 2019-09-30 | 24.364769 | 181.645401 | 175.044476 | 54.405307 |

WeChat: cstutorcs

Basic Analytics5

- `pandas` is quite error tolerant, in the sense that it captures errors and just puts a `NaN` value where the respective mathematical operation fails.
- Not only this, but as briefly shown before, one can also work with such incomplete data sets as if they were complete in a number of cases.
- This comes in handy, since reality is characterized by incomplete data sets more often than one might wish.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

NUMPY UNIVERSAL FUNCTIONS

In general, one can apply NumPy universal functions to `pandas DataFrame` objects whenever they could be applied to an `ndarray` object containing the same type of data.

- A `numpy.array` is a function that returns a `numpy.ndarray`.
- `numpy.ndarray()` is a class, while `numpy.array()` is a method / function to create `ndarray`

Basic Visualisation1: Introduction

- Plotting of data is only one line of code away in general, once the data is stored in a `DataFrame` object (see Figure 5-1):

```
In [53]: from pylab import plt, mpl
         plt.style.use('seaborn')
         mpl.rcParams['font.family'] = 'serif'
         %matplotlib inline

In [54]: df.cumsum().plot(lw=2.0, figsize=(10, 6));
```

Assignment Project Exam Help

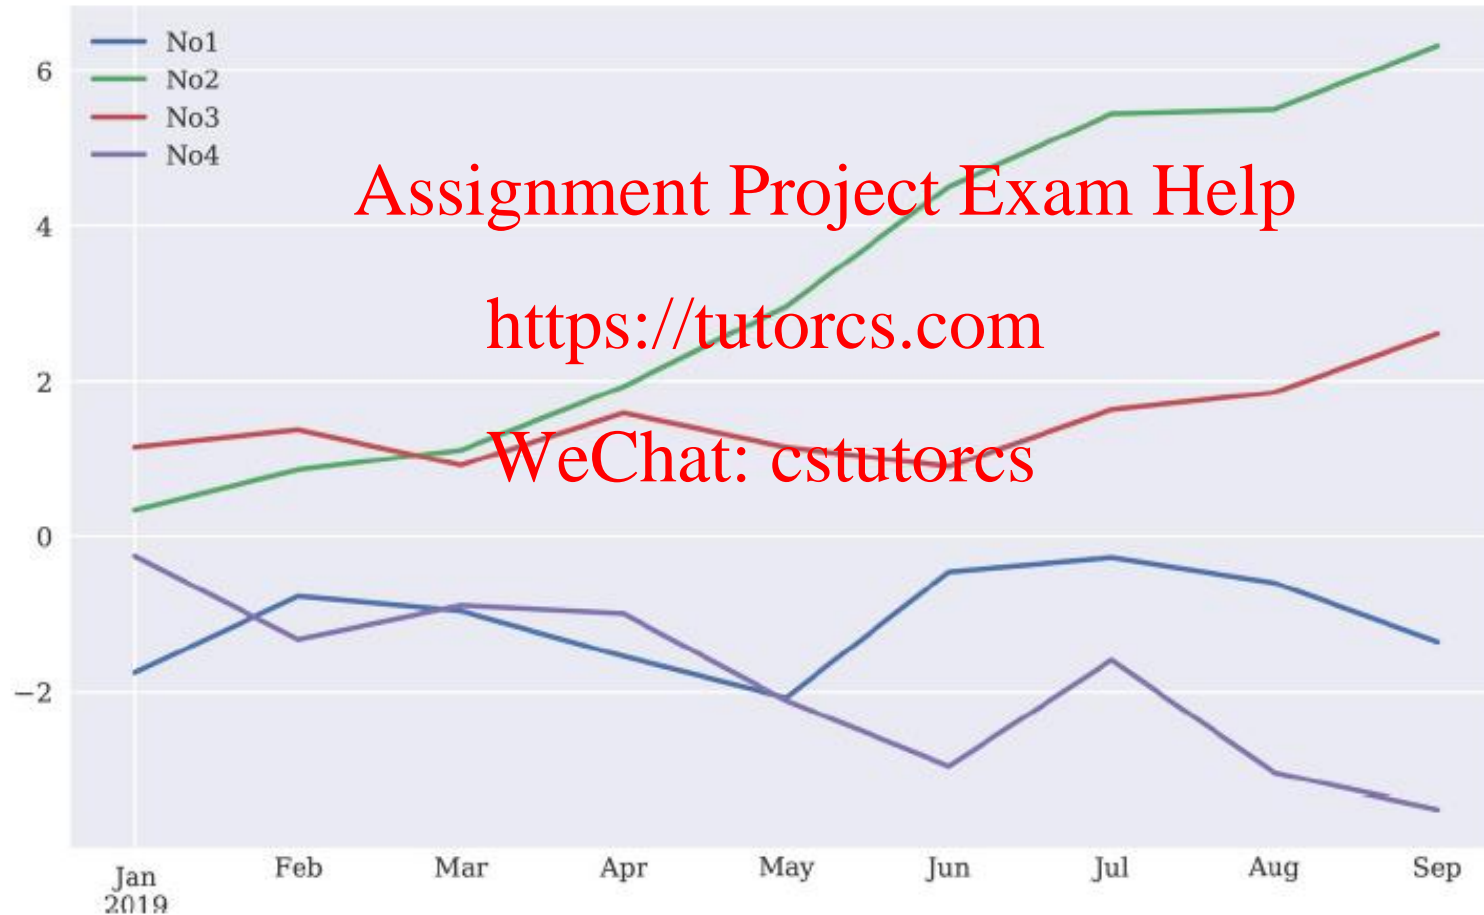
<https://tutorcs.com>

WeChat: cstutorcs

Plotting the cumulative sums of the four columns as a line plot.

- Basically, `pandas` provides a wrapper around `matplotlib` (see Chapter 7), specifically designed for `DataFrame` objects.
- Table 5-4 lists the parameters that the `plot()` method takes.

Basic Visualisation2: Line plot example



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Figure 5-1. Line plot of a DataFrame object

Basic Visualisation3:

Plot methods

Table 5-4. Parameters of plot() method

| Parameter | Format | Description |
|--------------|------------------------------|---|
| x | label/position, default None | Only used when column values are x-ticks |
| y | label/position, default None | Only used when column values are y-ticks |
| subplots | boolean, default False | Plot columns in subplots |
| sharex | boolean, default True | Share the x-axis |
| sharey | boolean, default False | Share the y-axis |
| use_index | boolean, default True | Use DataFrame.index as x-ticks |
| stacked | boolean, default False | Stack (only for bar plots) |
| sort_columns | boolean, default False | Sort columns alphabetically before plotting |

| Parameter | Format | Description |
|-------------|---|--|
| title | string, default None | Title for the plot |
| grid | boolean, default False | Show horizontal and vertical grid lines |
| legend | boolean, default True | Show legend of labels |
| ax | matplotlib axis object | matplotlib axis object to use for plotting |
| style | string or list/dictionary | Line plotting style (for each column) |
| kind | string (e.g. "line", "bar", "barh", "kde", "density") | Type of plot |
| logx | boolean, default False | Use logarithmic scaling of x-axis |
| logy | boolean, default False | Use logarithmic scaling of y-axis |
| xticks | sequence, default Index | X-ticks for the plot |
| yticks | sequence, default Values | Y-ticks for the plot |
| xlim | 2-tuple, list | Boundaries for x-axis |
| ylim | 2-tuple, list | Boundaries for y-axis |
| rot | integer, default None | Rotation of x-ticks |
| secondary_y | boolean/sequence, default False | Plot on secondary y-axis |
| mark_right | boolean, default True | Automatic labeling of secondary axis |
| colormap | string/colormap object, default None | Color map to use for plotting |
| kwds | keywords | Options to pass to matplotlib |

Basic Visualisation4: Bar plot example

- As another example, consider a bar plot of the same data (see Figure 5-2):

```
In [55]: df.plot.bar(figsize=(10, 6), rot=15);  
# df.plot(kind='bar', figsize=(10, 6))
```

❶

Plots the bar chart via `.plot.bar()`.

❷

Alternative syntax: uses the `kind` parameter to change the plot type.

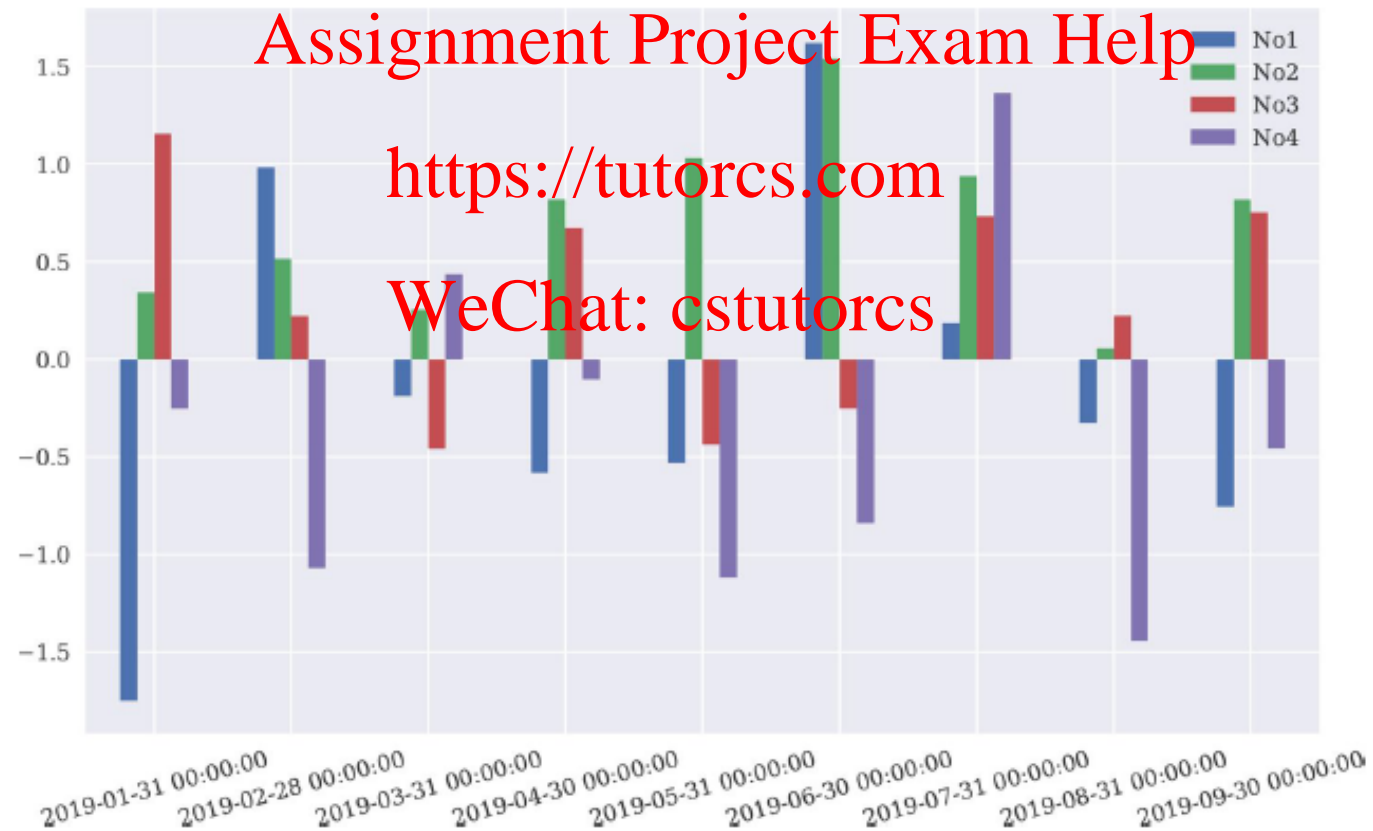


Figure 5-2. Bar plot of a DataFrame object

The Series Class

- So far this chapter has worked mainly with the pandas DataFrame class.
- Series is another important class that comes with pandas. It is characterized by the fact that it has only a single column of data.
- In that sense, it is a specialization of the DataFrame class that shares many but not all of its characteristics and capabilities.
- A Series object is obtained when a single column is selected from a multicolumn DataFrame object:

```
In [56]: type(df)
Out[56]: pandas.core.frame.DataFrame

In [57]: S = pd.Series(np.linspace(0, 15, 7), name='series')
```

```
In [58]: S
Out[58]: 0      0.0
         1      2.5
         2      5.0
         3      7.5
         4     10.0
         5     12.5
         6     15.0
         Name: series, dtype: float64
```

```
In [59]: type(S)
Out[59]: pandas.core.series.Series
```

```
In [60]: s = df['Nol']
Out[60]: s
Out[61]: 2019-01-31    -1.749765
         2019-02-28     0.981321
         2019-03-31    -0.189496
         2019-04-30    -0.583595
         2019-05-31    -0.531280
         2019-06-30     1.618982
         2019-07-31     0.184519
         2019-08-31    -0.326238
         2019-09-30    -0.756352
         Freq: M, Name: Nol, dtype: float64
```

```
In [62]: type(s)
Out[62]: pandas.core.series.Series
```

<https://tutorcs.com>

WeChat: cstutorcs

The Series Class2

- The main DataFrame methods are available for Series objects as well. For illustration, consider the mean() and plot() methods (see Figure 5-3):

```
In [63]: s.mean()  
Out[63]: -0.15021177307319458
```

```
In [64]: s.plot(lw=2.0, figsize=(10, 6));
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Figure 5-3. Line plot of a Series object

GroupBy Operations1

- pandas has powerful and flexible grouping capabilities.
- They work similarly to grouping in SQL as well as pivot tables in Microsoft Excel.
- To have something to group by one can add, for instance, a column indicating the quarter the respective data of the index belongs to:

Assignment Project Exam Help

<https://tutorcs.com>
WeChat: cstutorcs

```
In [65]: df['Quarter'] = ['Q1', 'Q1', 'Q1', 'Q2', 'Q2',  
                        'Q2', 'Q3', 'Q3', 'Q3']  
  
Out[65]:
```

| | No1 | No2 | No3 | No4 | Quarter |
|------------|-----------|----------|-----------|-----------|---------|
| 2019-01-31 | -1.749765 | 0.342680 | 1.153036 | -0.252436 | Q1 |
| 2019-02-28 | 0.981321 | 0.514219 | 0.221180 | -1.070043 | Q1 |
| 2019-03-31 | -0.189496 | 0.255001 | -0.458027 | 0.435163 | Q1 |
| 2019-04-30 | -0.583595 | 0.816847 | 0.672721 | -0.104411 | Q2 |
| 2019-05-31 | -0.531280 | 1.029733 | -0.438136 | -1.118318 | Q2 |
| 2019-06-30 | 1.618982 | 1.541605 | -0.251879 | -0.842436 | Q2 |
| 2019-07-31 | 0.184519 | 0.937082 | 0.731000 | 1.361556 | Q3 |
| 2019-08-31 | -0.326238 | 0.055676 | 0.222400 | -1.443217 | Q3 |
| 2019-09-30 | -0.756352 | 0.816454 | 0.750445 | -0.455947 | Q3 |

GroupBy Operations2

- The following code groups by the `Quarter` column and outputs statistics for the single groups:

```
In [66]: groups = df.groupby('Quarter') ❶
```

```
In [67]: groups.size() ❷
```

```
Out[67]: Quarter
Q1      3
Q2      3
Q3      3
dtype: int64
```

Assignment Project Exam Help

Groups according to the `Quarter` column.

```
In [68]: groups.mean() ❸
```

```
Out[68]:      No1      No2      No3      No4
Quarter
Q1    -0.319314  0.370634  0.305396 -0.295772
Q2     0.168035  1.129395 -0.005789 -0.668888
Q3    -0.299357  0.603071  0.567944 -0.219222
```

Gives the number of rows in each group.

<https://tutorcs.com>

WeChat: cstutorcs

Gives the mean per column.

```
In [69]: groups.max() ❹
```

```
Out[69]:      No1      No2      No3      No4
Quarter
Q1     0.981321  0.514219  1.153036  0.435163
Q2     1.618982  1.541605  0.672721 -0.104411
Q3     0.184519  0.937082  0.750445  1.361556
```

Gives the maximum value per column.

```
In [70]: groups.aggregate([min, max]).round(2) ❺
```

```
Out[70]:      No1      No2      No3      No4
Quarter
min  max  min  max  min  max  min  max
Q1   -1.75  0.98  0.26  0.51 -0.46  1.15 -1.07  0.44
Q2   -0.58  1.62  0.82  1.54 -0.44  0.67 -1.12 -0.10
Q3   -0.76  0.18  0.06  0.94  0.22  0.75 -1.44  1.36
```

Gives both the minimum and maximum values per column.

GroupBy Operations3

- Grouping can also be done with multiple columns.
- To this end, another column, indicating whether the month of the index date is odd or even, is introduced:

```
In [71]: df['Odd_Even'] = ['Odd', 'Even', 'Odd', 'Even', 'Odd', 'Even',  
                          'Odd', 'Even', 'Odd']  
  
In [72]: groups = df.groupby(['Quarter', 'Odd_Even'])  
  
In [73]: groups.size()  
Out[73]: Quarter Odd_Even  
Q1      Even      1  
        Odd       2  
Q2      Even      2  
        Odd       1  
Q3      Even      1  
        Odd       2  
dtype: int64  
  
In [74]: groups[['No1', 'No4']].aggregate([sum, np.mean])  
Out[74]:
```

| | | No1 | | No4 | |
|----|------|-----------|-----------|-----------|-----------|
| | | sum | mean | sum | mean |
| Q1 | Even | 0.981321 | 0.981321 | -1.070043 | -1.070043 |
| | Odd | -1.939261 | -0.969631 | 0.182727 | 0.091364 |
| Q2 | Even | 1.035387 | 0.517693 | -0.946847 | -0.473423 |
| | Odd | -0.531280 | -0.531280 | -1.118318 | -1.118318 |
| Q3 | Even | -0.326238 | -0.326238 | -1.443217 | -1.443217 |
| | Odd | -0.571834 | -0.285917 | 0.905609 | 0.452805 |

- This concludes the introduction to pandas and the use of DataFrame objects.

*Panel Data Pivot Table1

Create a spreadsheet-style pivot table as a DataFrame.

The levels in the pivot table will be stored in MultiIndex objects (hierarchical indexes) on the index and columns of the result DataFrame.

Function `pivot_table()`

- `pandas.pivot_table(data, values=None, index=None, columns=None, aggfunc='mean', fill_value=None, margins=False, dropna=True, margins_name='All', observed=False)`
- Create a spreadsheet-style pivot table as a DataFrame. The levels in the pivot table will be stored in MultiIndex objects (hierarchical indexes) on the index and columns of the result DataFrame

Function `pivot()`

- `DataFrame.pivot(self, index=None, columns=None, values=None)`
- Reshape data (produce a “pivot” table) based on column values. Uses unique values from specified *index* / *columns* to form axes of the resulting DataFrame. This function does not support data aggregation, multiple values will result in a MultiIndex in the columns.

Function `transpose()`

- `DataFrame.transpose(self, *args, **kwargs)`
- Reflect the DataFrame over its main diagonal by writing rows as columns and vice-versa.

Pivot

df

| | foo | bar | baz | zoo |
|---|-----|-----|-----|-----|
| 0 | one | A | 1 | x |
| 1 | one | B | 2 | y |
| 2 | one | C | 3 | z |
| 3 | two | A | 4 | q |
| 4 | two | B | 5 | w |
| 5 | two | C | 6 | t |



```
df.pivot(index='foo',  
         columns='bar',  
         values='baz')
```

| bar | A | B | C |
|-----|---|---|---|
| foo | | | |
| one | 1 | 2 | 3 |
| two | 4 | 5 | 6 |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Resources:

- https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.pivot_table.html
- <https://jakevdp.github.io/PythonDataScienceHandbook/03.09-pivot-tables.html>
- https://www.geeksforgeeks.org/python-pandas-pivot_table/
- <https://www.geeksforgeeks.org/python-pandas-dataframe-transpose/>

Complex Selection1

- Often, data selection is accomplished by formulation of conditions on column values, and potentially combining multiple such conditions logically.
- Consider the following data set:

```
In [75]: data = np.random.standard_normal((10, 2)) ❶
```

```
In [76]: df = pd.DataFrame(data, columns=['x', 'y']) ❷
```

```
In [77]: df.info() ❷
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 2 columns):
x      10 non-null float64
y      10 non-null float64
dtypes: float64(2)
memory usage: 240.0 bytes
```

```
In [78]: df.head() ❸
Out[78]:
```

| | x | y |
|---|-----------|-----------|
| 0 | 1.189622 | -1.690617 |
| 1 | -1.356399 | -1.232435 |
| 2 | -0.544439 | -0.668172 |
| 3 | 0.007315 | -0.612939 |
| 4 | 1.299748 | -1.733096 |

```
In [79]: df.tail() ❹
Out[79]:
```

| | x | y |
|---|-----------|-----------|
| 5 | -0.983310 | 0.357508 |
| 6 | -1.613579 | 1.470714 |
| 7 | -1.188018 | -0.549746 |
| 8 | -0.940046 | -0.827932 |
| 9 | 0.108863 | 0.507810 |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

❶ ndarray object with standard normally distributed random numbers.

❷ DataFrame object with the same random numbers.

❸ The first five rows via the head() method.

❹ The final five rows via the tail() method.

Complex Selection2

- The following code illustrates the application of Python's comparison operators and logical operators on values in the two columns:

- ① Check whether value in column `x` is greater than 0.5
- ② Check whether value in column `x` is positive *and* value in column `y` is negative.
- ③ Check whether value in column `x` is positive *or* value in column `y` is negative.

```
In [80]: df['x'] > 0.5 ①
Out[80]: 0      True
         1     False
         2     False
         3     False
         4      True
         5     False
         6     False
         7     False
         8     False
         9     False
         Name: x, dtype: bool
```

```
In [81]: (df['x'] > 0) & (df['y'] < 0) ②
Out[81]: 0      True
         1     False
         2     False
         3      True
         4      True
         5     False
         6     False
         7     False
         8     False
         9     False
         dtype: bool
```

```
In [82]: (df['x'] > 0) | (df['y'] < 0) ③
Out[82]: 0      True
         1      True
         2      True
         3      True
         4      True
         5     False
         6     False
         7      True
         8      True
         9      True
         dtype: bool
```

Complex Selection3

- Using the resulting Boolean Series objects, complex data (row) selection is straightforward.
- Alternatively, one can use the `query()` method and pass the conditions as `str` objects:

①

All rows for which the value in column `x` is greater than 0.5. *Note bool type

②

All rows for which the value in column `x` is positive *and* the value in column `y` is negative.

③

All rows for which the value in column `x` is positive *or* the value in column `y` is negative (columns are accessed here via the respective attributes).

```
In [83]: df[df['x'] > 0] ①
Out[83]:
```

| | x | y |
|---|----------|-----------|
| 0 | 1.189622 | -1.690617 |
| 3 | 0.007315 | -0.612939 |
| 4 | 1.299748 | -1.733096 |
| 9 | 0.108863 | 0.507810 |

```
In [84]: df.query('x > 0') ①
Out[84]:
```

| | x | y |
|---|----------|-----------|
| 0 | 1.189622 | -1.690617 |
| 3 | 0.007315 | -0.612939 |
| 4 | 1.299748 | -1.733096 |
| 9 | 0.108863 | 0.507810 |

```
In [85]: df[(df['x'] > 0) & (df['y'] < 0)] ②
Out[85]:
```

| | x | y |
|---|----------|-----------|
| 0 | 1.189622 | -1.690617 |
| 3 | 0.007315 | -0.612939 |
| 4 | 1.299748 | -1.733096 |

```
In [86]: df.query('x > 0 & y < 0') ②
Out[86]:
```

| | x | y |
|---|----------|-----------|
| 0 | 1.189622 | -1.690617 |
| 3 | 0.007315 | -0.612939 |
| 4 | 1.299748 | -1.733096 |

```
In [87]: df[(df.x > 0) | (df.y < 0)] ③
Out[87]:
```

| | x | y |
|---|-----------|-----------|
| 0 | 1.189622 | -1.690617 |
| 1 | -1.356399 | -1.232435 |
| 2 | -0.544439 | -0.668172 |
| 3 | 0.007315 | -0.612939 |
| 4 | 1.299748 | -1.733096 |
| 7 | -1.188018 | -0.549746 |
| 8 | -0.940046 | -0.827932 |
| 9 | 0.108863 | 0.507810 |

Complex Selection4

- Comparison operators can also be applied to complete DataFrame objects at once:

```
In [88]: df > 0 ❶
Out[88]:
```

| | x | y |
|---|-------|-------|
| 0 | True | False |
| 1 | False | False |
| 2 | False | False |
| 3 | True | False |
| 4 | True | False |
| 5 | False | True |
| 6 | False | True |
| 7 | False | False |
| 8 | False | False |
| 9 | True | True |

Assignment Project Exam Help

❶ Which values in the DataFrame object are positive?
<https://tutorcs.com>

❷ Select all such values and put a NaN in all other places.
WeChat: cstutorcs

```
In [89]: df[df > 0] ❷
Out[89]:
```

| | x | y |
|---|----------|----------|
| 0 | 1.189622 | NaN |
| 1 | NaN | NaN |
| 2 | NaN | NaN |
| 3 | 0.007315 | NaN |
| 4 | 1.299748 | NaN |
| 5 | NaN | 0.357508 |
| 6 | NaN | 1.470714 |
| 7 | NaN | NaN |
| 8 | NaN | NaN |
| 9 | 0.108863 | 0.507810 |

Concatenation, Joining and Merging

- This section walks through different approaches to combine two simple data sets in the form of DataFrame objects. The two simple data sets are:

```
In [90]: df1 = pd.DataFrame(['100', '200', '300', '400'],  
                             index=['a', 'b', 'c', 'd'],  
                             columns=['A',])
```

```
In [91]: df1  
Out[91]:
```

| | A |
|---|-----|
| a | 100 |
| b | 200 |
| c | 300 |
| d | 400 |

```
In [92]: df2 = pd.DataFrame(['200', '150', '50'],  
                             index=['f', 'b', 'd'],  
                             columns=['B',])
```

```
In [93]: df2  
Out[93]:
```

| | B |
|---|-----|
| f | 200 |
| b | 150 |
| d | 50 |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Concatenation

- *Concatenation* or *appending* basically means that rows are added from one DataFrame object to another one.
- This can be accomplished via the `append()` method or via the `pd.concat()` function.
- A major consideration is how the index values are handled:

```
In [94]: df1.append(df2, sort=False) ❶
```

```
Out[94]:
```

| | A | B |
|---|-----|-----|
| a | 100 | NaN |
| b | 200 | NaN |
| c | 300 | NaN |
| d | 400 | NaN |
| f | NaN | 200 |
| b | NaN | 150 |
| d | NaN | 50 |

```
In [95]: df1.append(df2, ignore_index=True, sort=False) ❷
```

```
Out[95]:
```

| | A | B |
|---|-----|-----|
| 0 | 100 | NaN |
| 1 | 200 | NaN |
| 2 | 300 | NaN |
| 3 | 400 | NaN |
| 4 | NaN | 200 |
| 5 | NaN | 150 |
| 6 | NaN | 50 |

```
In [96]: pd.concat((df1, df2), sort=False) ❸
```

```
Out[96]:
```

| | A | B |
|---|-----|-----|
| a | 100 | NaN |
| b | 200 | NaN |
| c | 300 | NaN |
| d | 400 | NaN |
| f | NaN | 200 |
| b | NaN | 150 |
| d | NaN | 50 |

```
In [97]: pd.concat((df1, df2), ignore_index=True, sort=False) ❹
```

```
Out[97]:
```

| | A | B |
|---|-----|-----|
| 0 | 100 | NaN |
| 1 | 200 | NaN |
| 2 | 300 | NaN |
| 3 | 400 | NaN |
| 4 | NaN | 200 |
| 5 | NaN | 150 |
| 6 | NaN | 50 |

❶ Appends data from `df2` to `df1` as new rows

❷ Does the same but ignores the indices.

❸ Has the same effect as the first append operation.

❹ Has the same effect as the second append operation.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Joining1

- When joining the two data sets, the sequence of the DataFrame objects also matters but in a different way.
- Only the index values from the first DataFrame object are used.
- This default behaviour is called a *left join*:

```
In [98]: df1.join(df2)
```

```
Out[98]:
```

| | A | B |
|---|-----|-----|
| a | 100 | NaN |
| b | 200 | 150 |
| c | 300 | NaN |
| d | 400 | 50 |

```
In [99]: df2.join(df1)
```

```
Out[99]:
```

| | B | A |
|---|-----|-----|
| f | 200 | NaN |
| b | 150 | 200 |
| d | 50 | 400 |

① Assignment Project Exam Help

① <https://tutorcs.com>

Index values of df1 are relevant.

② WeChat: cstutorcs

Index values of df2 are relevant.

Joining2

Documentation:

<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.join.html?highlight=join#pandas.DataFrame.join>

- There are a total of four different join methods available, each leading to a different behaviour with regard to how index values and the corresponding data rows are handled:

```
In [100]: df1.join(df2, how='left')  
Out[100]:
```

| | A | B |
|---|-----|-----|
| a | 100 | NaN |
| b | 200 | 150 |
| c | 300 | NaN |
| d | 400 | 50 |

Left join is the default operation.

```
In [101]: df1.join(df2, how='right')  
Out[101]:
```

| | A | B |
|---|-----|-----|
| f | NaN | 200 |
| b | 200 | 150 |
| d | 400 | 50 |

Right join is the same as reversing the sequence of the DataFrame objects.

```
In [102]: df1.join(df2, how='inner')  
Out[102]:
```

| | A | B |
|---|-----|-----|
| b | 200 | 150 |
| d | 400 | 50 |

Inner join only preserves those index values found in both indices.

```
In [103]: df1.join(df2, how='outer')  
Out[103]:
```

| | A | B |
|---|-----|-----|
| a | 100 | NaN |
| b | 200 | 150 |
| c | 300 | NaN |
| d | 400 | 50 |
| f | NaN | 200 |

Outer join preserves all index values from both indices.

how : {'left', 'right', 'outer', 'inner'}, default 'left'

How to handle the operation of the two objects.

- left: use calling frame's index (or column if on is specified)
- right: use *other's* index.
- outer: form union of calling frame's index (or column if on is specified) with *other's* index, and sort it. lexicographically.
- inner: form intersection of calling frame's index (or column if on is specified) with *other's* index, preserving the order of the calling's one.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Joining3

- A join can also happen based on an empty DataFrame object.
- In this case, the columns are created *sequentially*, leading to behaviour similar to a left join:

```
In [104]: df = pd.DataFrame()
```

```
In [105]: df['A'] = df1['A']
```

```
In [106]: df
```

```
Out[106]:
```

| | A |
|---|-----|
| a | 100 |
| b | 200 |
| c | 300 |
| d | 400 |

```
In [107]: df['B'] = df2
```

```
In [108]: df
```

```
Out[108]:
```

| | A | B |
|---|-----|-----|
| a | 100 | NaN |
| b | 200 | 150 |
| c | 300 | NaN |
| d | 400 | 50 |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Joining4

- Making use of a dictionary to combine the data sets yields a result similar to an outer join since the columns are created *simultaneously*:

```
In [109]: df = pd.DataFrame({'A': df1['A'], 'B': df2['B']})
```

```
In [110]: df
```

```
Out[110]:
```

| | A | B |
|---|-----|-----|
| a | 100 | NaN |
| b | 200 | 150 |
| c | 300 | NaN |
| d | 400 | 50 |
| f | NaN | 200 |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



The columns of the DataFrame objects are used as values in the dict object.

Merging1

- While a join operation takes place based on the indices of the DataFrame objects to be joined, a merge operation typically takes place on a column shared between the two data sets.
- To this end, a new column C is added to both original DataFrame objects:

```
In [111]: c = pd.Series([250, 150, 50, 100], index=['a', 'b', 'c', 'd'], 'c')
          df1['C'] = c
          df2['C'] = c
```

```
In [112]: df1
Out[112]:
```

| | A | C |
|---|-----|-------|
| a | 100 | NaN |
| b | 200 | 250.0 |
| c | 300 | 50.0 |
| d | 400 | 150.0 |

```
In [113]: df2
Out[113]:
```

| | B | C |
|---|-----|-------|
| f | 200 | NaN |
| b | 150 | 250.0 |
| d | 50 | 150.0 |

Assignment Project Exam Help

<https://tutores.com>

WeChat: cstutorcs

Merging2

- By default, the merge operation in this case takes place based on the single shared column C.
- Other options are available, however, such as an *outer* merge:

Assignment Project Exam Help

In [114]: `pd.merge(df1, df2)` ❶

Out[114]:

| | A | C | B |
|---|-----|-------|-----|
| 0 | 100 | NaN | 200 |
| 1 | 200 | 250.0 | 150 |
| 2 | 400 | 150.0 | 50 |

<https://tutorcs.com>

WeChat: cstutorcs

❶

The default merge on column c.

❷

In [115]: `pd.merge(df1, df2, on='C')` ❶

Out[115]:

| | A | C | B |
|---|-----|-------|-----|
| 0 | 100 | NaN | 200 |
| 1 | 200 | 250.0 | 150 |
| 2 | 400 | 150.0 | 50 |

An outer merge is also possible, preserving all data rows.

In [116]: `pd.merge(df1, df2, how='outer')` ❷

Out[116]:

| | A | C | B |
|---|-----|-------|-----|
| 0 | 100 | NaN | 200 |
| 1 | 200 | 250.0 | 150 |
| 2 | 300 | 50.0 | NaN |
| 3 | 400 | 150.0 | 50 |

Merging3

- Many more types of merge operations are available, a few of which are illustrated in the following code:

left_on : *label or list, or array-like*

Column or index level names to join on in the left DataFrame. Can also be an array or list of arrays of the length of the left DataFrame. These arrays are treated as if they are columns.

right_on : *label or list, or array-like*

Column or index level names to join on in the right DataFrame. Can also be an array or list of arrays of the length of the right DataFrame. These arrays are treated as if they are columns.

```
In [117]: pd.merge(df1, df2, left_on='A', right_on='B')
Out[117]:
```

| | A | C_x | B | C_y |
|---|-----|-------|-----|-----|
| 0 | 200 | 250.0 | 200 | NaN |

```
In [118]: pd.merge(df1, df2, left_on='A', right_on='B', how='outer')
Out[118]:
```

| | A | C_x | B | C_y |
|---|-----|-------|-----|-------|
| 0 | 100 | NaN | NaN | NaN |
| 1 | 200 | 250.0 | 200 | NaN |
| 2 | 300 | 50.0 | NaN | NaN |
| 3 | 400 | 150.0 | NaN | NaN |
| 4 | NaN | NaN | 150 | 250.0 |
| 5 | NaN | NaN | 50 | 150.0 |

```
In [119]: pd.merge(df1, df2, left_index=True, right_index=True)
Out[119]:
```

| | A | C_x | B | C_y |
|---|-----|-------|-----|-------|
| b | 200 | 250.0 | 150 | 250.0 |
| d | 400 | 150.0 | 50 | 150.0 |

```
In [120]: pd.merge(df1, df2, on='C', left_index=True)
Out[120]:
```

| | A | C | B |
|---|-----|-------|-----|
| f | 100 | NaN | 200 |
| b | 200 | 250.0 | 150 |
| d | 400 | 150.0 | 50 |

```
In [121]: pd.merge(df1, df2, on='C', right_index=True)
Out[121]:
```

| | A | C | B |
|---|-----|-------|-----|
| a | 100 | NaN | 200 |
| b | 200 | 250.0 | 150 |
| d | 400 | 150.0 | 50 |

```
In [122]: pd.merge(df1, df2, on='C', left_index=True, right_index=True)
Out[122]:
```

| | A | C | B |
|---|-----|-------|-----|
| b | 200 | 250.0 | 150 |
| d | 400 | 150.0 | 50 |

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

* Merging Extra Example

3. Merging using `left_on` and `right_on`

It might happen that the column on which you want to merge the DataFrames have different names. For such merges, you will have to specify the `left_on` as the left DataFrame name and `right_on` as the right DataFrame name, for example:

```
pd.merge(
    df_customer,
    df_info_2,
    left_on='id',
    right_on='customer_id'
)
```

Extra example

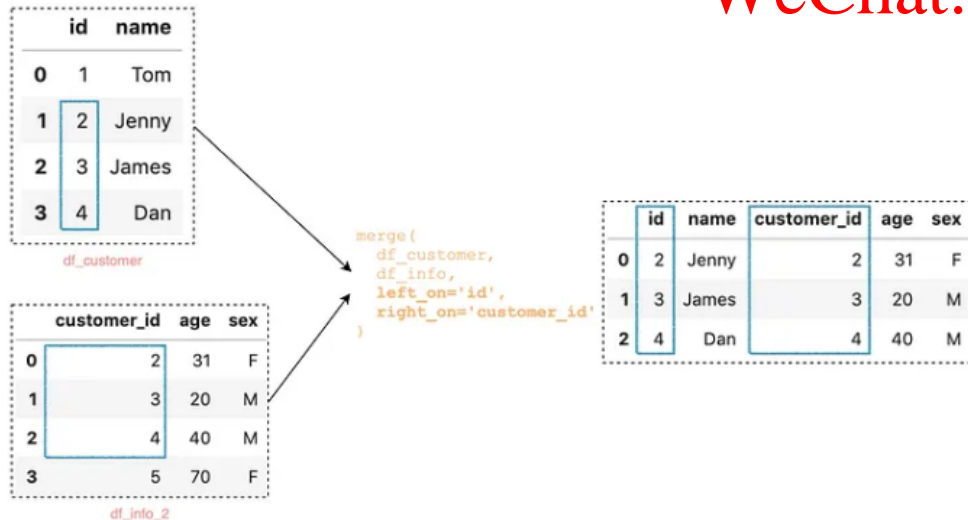
```
df_customer = pd.DataFrame({
    'id': [1, 2, 3, 4],
    'name': ['Tom', 'Jenny', 'James', 'Dan'],
})
df_info = pd.DataFrame({
    'customer_id': [2, 3, 4, 5],
    'age': [31, 20, 40, 70],
    'sex': ['F', 'M', 'M', 'F']
})
```

```
result = pd.merge(df_customer, df_info, left_on='id', right_on='customer_id')
print(result)
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



<https://towardsdatascience.com/all-the-pandas-merge-you-should-know-for-combining-datasets-526b9ecaf184>

Performance Aspects1

- Many examples in this chapter illustrate that there are often multiple options to achieve the same goal with pandas.
- This section compares such options for adding up two columns element-wise.
- First, the data set, generated with NumPy:

```
In [123]: data = np.random.standard_normal((1000000, 2))
```

```
In [124]: data.nbytes  
Out[124]: 16000000
```

```
In [125]: df = pd.DataFrame(data, columns=['x', 'y'])
```

```
In [126]: df.info()  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1000000 entries, 0 to 999999  
Data columns (total 2 columns):  
x      1000000 non-null float64  
y      1000000 non-null float64  
dtypes: float64(2)  
memory usage: 15.3 MB
```

The ndarray object with random numbers.

The DataFrame object with the random numbers.

Performance Aspects2

- Second, some options to accomplish the task at hand with performance values:

```
In [127]: %time res = df['x'] + df['y'] ❶
CPU times: user 7.35 ms, sys: 7.43 ms, total: 14.8 ms
Wall time: 7.48 ms
```

```
In [128]: res[:3]
Out[128]: 0    0.387242
          1   -0.969343
          2   -0.863159
          dtype: float64
```

```
In [129]: %time res = df.sum(axis=1) ❷
CPU times: user 130 ms, sys: 30.6 ms, total: 161 ms
Wall time: 101 ms
```

```
In [130]: res[:3]
Out[130]: 0    0.387242
          1   -0.969343
          2   -0.863159
          dtype: float64
```

```
In [131]: %time res = df.values.sum(axis=1) ❸
CPU times: user 50.3 ms, sys: 2.75 ms, total: 53.1 ms
Wall time: 27.9 ms
```

```
In [132]: res[:3]
Out[132]: array([ 0.3872424 , -0.96934273, -0.86315944])
```

```
In [133]: %time res = np.sum(df, axis=1) ❹
CPU times: user 127 ms, sys: 15.1 ms, total: 142 ms
Wall time: 73.7 ms
```

```
In [134]: res[:3]
Out[134]: 0    0.387242
          1   -0.969343
          2   -0.863159
          dtype: float64
```

```
In [135]: %time res = np.sum(df.values, axis=1) ❺
CPU times: user 49.3 ms, sys: 2.36 ms, total: 51.7 ms
Wall time: 26.9 ms
```

```
In [136]: res[:3]
Out[136]: array([ 0.3872424 , -0.96934273, -0.86315944])
```

❶ Working with the columns (Series objects) directly is the fastest approach.

❷ This calculates the sums by calling the `sum()` method on the DataFrame object.

❸ This calculates the sums by calling the `sum()` method on the ndarray object.

❹ This calculates the sums by using the function `np.sum()` on the DataFrame object.

❺ This calculates the sums by using the function `np.sum()` on the ndarray object.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutores

Performance Aspects3

- Finally, two more options which are based on the methods `eval()` and `apply()`, respectively:1

```
In [137]: %time res = df.eval('x + y')
CPU times: user 25.5 ms, sys: 17.2 ms, total: 42.7 ms
Wall time: 22.5 ms
```

①

`eval()` is a method dedicated to evaluation of (complex) numerical expressions; columns can be directly addressed.

```
In [138]: res[:3]
Out[138]: 0    0.387242
          1   -0.969343
          2   -0.863159
          dtype: float64
```

②

The slowest option is to use the `apply()` method row-by-row; this is like looping on the Python level over all rows.

```
In [139]: %time res = df.apply(lambda row: row['x'] + row['y'], axis=1)
CPU times: user 19.6 s, sys: 83.3 ms, total: 19.7 s
Wall time: 19.9 s
```

②

```
In [140]: res[:3]
Out[140]: 0    0.387242
          1   -0.969343
          2   -0.863159
          dtype: float64
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

CHOOSE WISELY

pandas often provides multiple options to accomplish the same goal. If unsure of which to use, compare the options to verify that the best possible performance is achieved when time is critical. In this simple example, execution times differ by orders of magnitude.

Conclusion

- `pandas` is a powerful tool for data analysis and has become the central package in the so-called *PyData* stack.
- Its `DataFrame` class is particularly suited to working with tabular data of any kind.
- Most operations on such objects are vectorized, leading not only — as in the `NumPy` case — to concise code but also to high performance in general.
- In addition, `pandas` makes working with incomplete data sets convenient (which is not the case with `NumPy`, for instance).
- `pandas` and the `DataFrame` class will be central in many later chapters of the book, where additional features will be used and introduced when necessary.
- **Please also consult Wes McKinney – Python for Data Analysis 2018 2nd Edition Chapter 8 Data Wrangling, Join and Reshape & Chapter 10 Data Aggregation and Group Operations for further knowledge acquisition on these subjects.**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs