

Assignment Project Exam Help

5QQMN534: Algorithmic Finance

<https://tutorcs.com>
WeChat: cstutorcs

Week6: Financial Data Extraction and Time Series Analysis

Yves Hilpisch - Python for Finance 2nd Edition 2019: Chapter 8

Agenda

- Financial Data
 - Data Import
 - Summary Statistics
 - Changes over Time
 - Resampling
- Rolling Statistics
 - An overview
 - Technical Analysis Example
- Correlation Analysis
 - The Data
 - Logarithmic Returns
 - Correlation
- High Frequency Data
- Conclusion

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Financial Data

- Financial time series data is one of the most important types of data in finance.
- This is data indexed by date and/or time.
- For example, prices of stocks over time represent financial time series data.
- Similarly, the EUR/USD exchange rate over time represents a financial time series; the exchange rate is quoted in brief intervals of time, and a collection of such quotes then is a time series of exchange rates.
- There is no financial discipline that gets by without considering time an important factor.
- This mainly is the same as with physics and other sciences.
- The major tool to cope with time series data in Python is pandas.
- Wes McKinney, the original and main author of pandas, started developing the library when working as an analyst at AQR Capital Management, a large hedge fund.
- It is safe to say that pandas has been designed from the ground up to work with financial time series data.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Financial Data

The chapter is mainly based on two financial time series data sets in the form of comma-separated values (CSV) files. It proceeds along the following lines:

“Financial Data”

This section is about the basics of working with financial times series data using `pandas`: data import, deriving summary statistics, calculating changes over time, and resampling.

“Rolling Statistics”

In financial analysis, rolling statistics play an important role. These are statistics calculated in general over a fixed time interval that is *rolled forward* over the complete data set. A popular example is simple moving averages. This section illustrates how `pandas` supports the calculation of such statistics.

“Correlation Analysis”

This section presents a case study based on financial time series data for the S&P 500 stock index and the VIX volatility index. It provides some support for the stylized (empirical) fact that both indices are negatively correlated.

“High-Frequency Data”

This section works with high-frequency data, or *tick data*, which has become commonplace in finance. `pandas` again proves powerful in handling such data sets.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Financial Data

- This section works with a locally stored financial data set in the form of a CSV file. Technically, such files are simply text files with a data row structure characterized by commas that separate single values. Before importing the data, some package imports and customizations:

Assignment Project Exam Help

<https://tutorcs.com>

```
In [1]: import numpy as np
import pandas as pd
from pylab import mpl, plt
plt.style.use('seaborn')
mpl.rcParams['font.family'] = 'serif'
%matplotlib inline
```

WeChat: estutorcs

Data Import1

- pandas provides a number of different functions and DataFrame methods to
- import data stored in different formats (CSV, SQL, Excel, etc.) and to export data to different formats (see Chapter 9 for more details).
- The following code uses the `pd.read_csv()` function to import the time series data set from the CSV file:

```
In [2]: filename = '../..source/tr_eikon_eod_data.csv' ❶

In [3]: f = open(filename, 'r') ❷
f.readlines()[:5] ❷
Out[3]: ['Date,AAPL.O,MSFT.O,INTC.O,AMZN.O,GS.N,SPY,.SPX,.VIX,EUR=,XAU=,GDX,
,GLD\n',
'2010-01-01,,,,,,,,,1.4323,1096.35,,\n',
'2010-01-04,30.57282657,30.95,20.88,133.9,173.08,113.33,1132.99,20.04,
,1.4411,1120.0,47.71,109.8\n',
'2010-01-05,30.625683660000004,30.96,20.87,134.69,176.14,113.63,1136.52,
,19.35,1.4368,1118.65,48.17,109.7\n',
'2010-01-06,30.138541290000003,30.77,20.8,132.25,174.26,113.71,1137.14,
,19.16,1.4412,1138.5,49.34,111.51\n']

In [4]: data = pd.read_csv(filename, ❸
                                index_col=0, ❹
                                parse_dates=True) ❺

In [5]: data.info() ❻
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2216 entries, 2010-01-01 to 2018-06-29
Data columns (total 12 columns):
AAPL.O      2138 non-null float64
MSFT.O      2138 non-null float64
INTC.O      2138 non-null float64
AMZN.O      2138 non-null float64
GS.N        2138 non-null float64
SPY         2138 non-null float64
.SPX        2138 non-null float64
.VIX        2138 non-null float64
EUR=        2216 non-null float64
XAU=        2211 non-null float64
GDX         2138 non-null float64
GLD         2138 non-null float64
dtypes: float64(12)
memory usage: 225.1 KB
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Note: The filename in code has been set to a directory path.

The file contains end-of-day (EOD) data for different financial instruments as retrieved from the Thomson Reuters Eikon Data API.

- ❶ Specifies the path and filename.
- ❷ Shows the first five rows of the raw data (Linux/Mac).
- ❸ The filename passed to the `pd.read_csv()` function.
- ❹ Specifies that the first column shall be handled as an index.
- ❺ Specifies that the index values are of type `datetime`.
- ❻ The resulting `DataFrame` object.

Data Import2

- At this stage, a financial analyst probably takes a first look at the data, either by inspecting or visualizing it (see Figure 8-1) (next slide):

```
In [6]: data.head() ❶
Out[6]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX
Date								
2010-01-01	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2010-01-04	30.572827	30.950	20.88	133.90	173.08	113.33	1132.99	20.04
2010-01-05	30.625684	30.960	20.87	134.69	176.14	113.63	1136.52	19.35
2010-01-06	30.138541	30.770	20.80	132.25	174.26	113.71	1137.14	19.16
2010-01-07	30.082827	30.452	20.60	130.00	177.67	114.19	1141.69	19.06

```
EUR= XAU= GDX GLD
Date
2010-01-01 1.4328 1096.35 NaN NaN
2010-01-04 1.4412 1120.00 49.71 109.80
2010-01-05 1.4368 1118.65 48.17 109.70
2010-01-06 1.4412 1138.50 49.34 111.51
2010-01-07 1.4318 1131.90 49.10 110.82
```

```
In [7]: data.tail() ❷
Out[7]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX
Date								
2018-06-25	182.17	98.39	50.71	1663.15	221.54	271.00	2717.07	17.33
2018-06-26	184.43	99.08	49.67	1691.09	221.58	271.60	2723.06	15.92
2018-06-27	184.16	97.54	48.76	1660.51	220.18	269.35	2699.63	17.91
2018-06-28	185.50	98.63	49.25	1701.45	223.42	270.89	2716.31	16.85
2018-06-29	185.11	98.61	49.71	1699.80	220.57	271.28	2718.37	16.09

```
EUR= XAU= GDX GLD
Date
2018-06-25 1.1702 1265.00 22.01 119.89
2018-06-26 1.1645 1258.64 21.95 119.26
2018-06-27 1.1552 1251.62 21.81 118.58
2018-06-28 1.1567 1247.88 21.93 118.22
2018-06-29 1.1683 1252.25 22.31 118.65
```

```
In [8]: data.plot(figsize=(10, 12), subplots=True); ❸
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Data Import3

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs




Figure 8-1. Financial time series data as line plots

Data Import4

- The data used is from the Thomson Reuters (TR) Eikon Data API. In the TR world symbols for financial instruments are called *Reuters Instrument Codes* (RICs). The financial instruments that the single RICs represent are:

```
In [9]: instruments = ['Apple Stock', 'Microsoft Stock',  
                      'Intel Stock', 'Amazon Stock', 'Goldman Sachs Stock',  
                      'SPDR S&P 500 ETF Trust', 'S&P 500 Index',  
                      'VIX Volatility Index', 'EUR/USD Exchange Rate',  
                      'Gold Price', 'VanEck Vectors Gold Miners ETF',  
                      'SPDR Gold Trust']
```

```
In [10]: for ric, name in zip(data.columns, instruments):  
          print('{:8s} | {}'.format(ric, name))  
AAPL.O   | Apple Stock  
MSFT.O   | Microsoft Stock  
INTC.O   | Intel Stock  
AMZN.O   | Amazon Stock  
GS.N     | Goldman Sachs Stock  
SPY      | SPDR S&P 500 ETF Trust  
.SPX     | S&P 500 Index  
.VIX     | VIX Volatility Index  
EUR=     | EUR/USD Exchange Rate  
XAU=     | Gold Price  
GDX      | VanEck Vectors Gold Miners ETF  
GLD      | SPDR Gold Trust
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

*Python zip built in function

The `zip()` function returns a zip object, which is an iterator of tuples where the first item in each passed iterator is paired together, and then the second item in each passed iterator are paired together etc. If the passed iterators have different lengths, the iterator with the least items decides the length of the new iterator.

Assignment Project Exam Help

https://www.w3schools.com/python/ref_func_zip.asp

<https://tutorcs.com>

WeChat: cstutorcs

```
a = ("John", "Charles", "Mike")  
b = ("Jenny", "Christy", "Monica")
```

```
x = zip(a, b)
```

```
#use the tuple() function to display a readable version of the result:
```

```
print(tuple(x))
```

```
(( 'John', 'Jenny'), ('Charles', 'Christy'), ('Mike', 'Monica'))
```

Summary Statistics

- The next step the financial analyst might take is to have a look at different summary statistics for the data set to get a “feeling” for what it is all about:

①

`info()` gives some metainformation about the `DataFrame` object.

②

`describe()` provides useful standard statistics per column.

QUICK INSIGHTS

`pandas` provides a number of methods to gain a quick overview over newly imported financial time series data sets, such as `info()` and `describe()`. They also allow for quick checks of whether the importing procedure worked as desired (e.g., whether the `DataFrame` object indeed has an index of type `DatetimeIndex`).

```
In [11]: data.info()
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2216 entries, 2010-01-01 to 2018-06-29
Data columns (total 12 columns):
AAPL.O      2138 non-null float64
MSFT.O      2138 non-null float64
INTC.O      2138 non-null float64
AMZN.O      2138 non-null float64
GS.N        2138 non-null float64
SPY         2138 non-null float64
.SPX        2138 non-null float64
.VIX        2138 non-null float64
EUR=        2216 non-null float64
XAU=        2211 non-null float64
GDX         2138 non-null float64
GLD         2138 non-null float64
dtypes: float64(12)
memory usage: 225.1 KB
```

```
In [12]: data.describe().round(2)
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX
count	2138.00	2138.00	2138.00	2138.00	2138.00	2138.00	2138.00	2138.00
mean	33.46	44.56	29.36	480.46	170.22	180.32	1802.71	17.03
std	40.55	19.53	8.17	372.31	42.48	48.19	483.34	5.88
min	27.44	23.01	17.66	108.61	87.70	102.20	1022.58	9.14
25%	60.29	28.57	22.51	213.60	146.61	133.99	1338.57	13.07
50%	90.55	39.66	27.33	322.06	164.43	186.32	1863.08	15.58
75%	117.24	54.37	34.71	698.85	192.13	210.99	2108.94	19.07
max	193.98	102.49	57.08	1750.08	273.38	286.58	2872.87	48.00

	EUR=	XAU=	GDX	GLD
count	2216.00	2211.00	2138.00	2138.00
mean	1.25	1349.01	33.57	130.09
std	0.11	188.75	15.17	18.78
min	1.04	1051.36	12.47	100.50
25%	1.13	1221.53	22.14	117.40
50%	1.27	1292.61	25.62	124.00
75%	1.35	1428.24	48.34	139.00
max	1.48	1898.99	66.63	184.59

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Summary Statistics

- There are also options, of course, to customize what types of statistic to derive and display:

```
In [13]: data.mean()
Out[13]: AAPL.O      93.455973
         MSFT.O      44.561115
         INTC.O      29.364192
         AMZN.O      480.461251
         GS.N        170.216221
         SPY         180.323029
         .SPX        1802.713106
         .VIX        17.027133
         EUR=        1.248587
         XAU=       1349.014130
         GDX         33.566525
         GLD         130.086590
         dtype: float64
```

```
In [14]: data.aggregate([min,
                          np.mean,
                          np.std,
                          np.median,
                          max])
Out[14]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUR=
min	93.45	23.01	17.66	108.61	87.70	102.20	1022.58	9.14	1.04
mean	93.46	44.56	29.36	480.46	170.22	180.32	1802.71	17.03	1.25
std	40.55	19.53	8.17	372.31	42.48	48.19	483.34	5.88	0.11
median	90.55	39.66	27.33	322.06	164.43	186.32	1863.08	15.58	1.27
max	193.98	102.49	57.08	1750.08	273.38	286.58	2872.87	48.00	1.48

	XAU=	GDX	GLD
min	1051.36	12.47	100.50
mean	1349.01	33.57	130.09
std	188.75	15.17	18.78
median	1292.61	25.62	124.00
max	1898.99	66.63	184.59

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutores

Using the `aggregate()` method also allows one to pass custom functions.

.aggregate()

pandas.DataFrame.aggregate

`DataFrame.aggregate(func=None, axis=0, *args, **kwargs)`

[\[source\]](#)

Returns:

scalar, Series or DataFrame

Aggregate using one or more operations over the specified axis.

Parameters: **func** : *function, str, list or dict*

Function to use for aggregating the data. If a function, must either work when passed a DataFrame or when passed to DataFrame.apply.

Accepted combinations are:

- function
- string function name
- list of functions and/or function names, e.g. `[np.sum, 'mean']`
- dict of axis labels -> functions, function names or list of such

axis : {0 or 'index', 1 or 'columns'}, default 0

If 0 or 'index': apply function to each column. If 1 or 'columns': apply function to each row.

***args**

Positional arguments to pass to *func*.

****kwargs**

Keyword arguments to pass to *func*.

- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.aggregate.html>
- Note by **default** axis = 0 applies to columns. See manual

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

The return can be:

- scalar : when Series.agg is called with single function
 - Series : when DataFrame.agg is called with a single function
 - DataFrame : when DataFrame.agg is called with several functions
- Return scalar, Series or DataFrame.

The aggregation operations are always performed over an axis, either the index (default) or the column axis. This behavior is different from *numpy* aggregation functions (*mean*, *median*, *prod*, *sum*, *std*, *var*), where the default is to compute the aggregation of the flattened array, e.g., `numpy.mean(arr_2d)` as opposed to `numpy.mean(arr_2d, axis=0)`.

agg is an alias for *aggregate*. Use the alias.

Changes over Time

- Statistical analysis methods are often based on changes over time and not the absolute values themselves.
- There are multiple options to calculate the changes in a time series over time, including absolute differences, percentage changes, and logarithmic (log) returns.
- First, the absolute differences, for which pandas provides a special method:

①

`diff()` provides the absolute changes between two index values.

②

Of course, aggregation operations can be applied in addition.

```
In [15]: data.diff().head()
Out[15]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUR=
Date									
2010-01-01	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2010-01-04	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.0088
2010-01-05	0.052857	0.010	-0.01	0.79	3.06	0.30	3.53	-0.69	-0.0043
2010-01-06	-0.487142	-0.190	-0.07	-2.44	-1.88	0.08	0.62	-0.19	0.0044
2010-01-07	-0.055714	-0.318	-0.20	-2.25	3.41	0.48	4.55	-0.10	-0.0094

	XAU=	GDX	GLD
Date			
2010-01-01	NaN	NaN	NaN
2010-01-04	23.65	NaN	NaN
2010-01-05	-1.35	0.46	-0.10
2010-01-06	19.85	1.17	1.81
2010-01-07	-5.60	-0.24	-0.69

```
In [16]: data.diff().mean()
```

```
Out[16]: AAPL.O    0.064737
          MSFT.O    0.031246
          INTC.O    0.013540
          AMZN.O    0.706608
          GS.N      0.028224
          SPY       0.072103
          .SPX      0.732659
          .VIX     -0.019583
          EUR=     -0.000119
          XAU=      0.041887
          GDX      -0.015071
          GLD      -0.003455
          dtype: float64
```

Absolute change refers to the simple difference in the indicator over two periods in time, i.e.

Absolute change = Value of indicator in period 2 – Value of indicator in period 1

.diff()

pandas.DataFrame.diff

`DataFrame.diff(periods=1, axis=0)`

First discrete difference of element.

Calculates the difference of a Dataframe element compared with another element in the Dataframe (default is element in previous row).

Parameters: **periods** : *int, default 1*

Periods to shift for calculating difference, accepts negative values.

axis : {0 or 'index', 1 or 'columns'}, *default 0*

Take difference over rows (0) or columns (1).

Returns: **Dataframe**

First differences of the Series.

Assignment Project Exam Help [source]

<https://tutorcs.com>

WeChat: cstutorcs

<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.diff.html>

Changes over Time2

- From a statistics point of view, absolute changes are not optimal because they are dependent on the scale of the time series data itself.
- Therefore, percentage changes are usually preferred.
- The following code derives the percentage changes or percentage returns (also: simple returns) in a financial context and visualizes their mean values per column (see **Figure 8-2**) (next slide):

```
In [17]: data.pct_change().round(3).head()
Out[17]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUR=
Date									
2010-01-04	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2010-01-05	0.002	0.000	-0.000	0.006	0.018	0.003	0.003	-0.034	-0.003
2010-01-06	-0.016	-0.006	-0.003	-0.018	-0.011	0.001	0.001	-0.010	0.003
2010-01-07	-0.002	-0.010	-0.010	-0.017	0.020	0.004	0.004	-0.005	-0.007

	XAU=	GDX	GLD
Date			
2010-01-04	NaN	NaN	NaN
2010-01-05	-0.001	0.010	-0.001
2010-01-06	0.018	0.024	0.016
2010-01-07	-0.006	-0.005	-0.006


```
In [18]: data.pct_change().mean().plot(kind='bar', figsize=(10, 6));
```

❶

`pct_change()` calculates the percentage change between two index values.

❷

The mean values of the results are visualized as a bar plot.

Changes over Time3




Figure 8-2. Mean values of percentage changes as bar plot

Colormap

<https://matplotlib.org/stable/tutorials/colors/colormaps.html>

supported values are 'Accent', 'Accent_r', 'Blues',
'Blues_r', 'BrBG', 'BrBG_r', 'BuGn', 'BuGn_r', 'BuPu',
'BuPu_r', 'CMRmap', 'CMRmap_r', 'Dark2', 'Dark2_r',
'GnBu', 'GnBu_r', 'Greens', 'Greens_r', 'Greys',
'Greys_r', 'OrRd', 'OrRd_r', 'Oranges', 'Oranges_r',
'PRGn', 'PRGn_r', 'Paired', 'Paired_r', 'Pastel1',
'Pastel1_r', 'Pastel2', 'Pastel2_r', 'PiYG', 'PiYG_r',
'PuBu', 'PuBuGn', 'PuBuGn_r', 'PuBu_r', 'PuOr',
'PuOr_r', 'PuRd', 'PuRd_r', 'Purples', 'Purples_r',
'RdBu', 'RdBu_r', 'RdGy', 'RdGy_r', 'RdPu', 'RdPu_r',
'RdYlBu', 'RdYlBu_r', 'RdYlGn', 'RdYlGn_r', 'Reds',
'Reds_r', 'Set1', 'Set1_r', 'Set2', 'Set2_r', 'Set3',
'Set3_r', 'Spectral', 'Spectral_r', 'Wistia',
'Wistia_r', 'YlGn', 'YlGnBu', 'YlGnBu_r', 'YlGn_r',
'YlOrBr', 'YlOrBr_r', 'YlOrRd', 'YlOrRd_r', 'afmhot',
'afmhot_r', 'autumn', 'autumn_r', 'binary',
'binary_r', 'bone', 'bone_r', 'brg', 'brg_r', 'bwr',
'bwr_r', 'cividis', 'cividis_r', 'cool', 'cool_r',

'coolwarm', 'coolwarm_r', 'copper', 'copper_r',
'cubehelix', 'cubehelix_r', 'flag', 'flag_r',
'gist_earth', 'gist_earth_r', 'gist_gray',
'gist_gray_r', 'gist_heat', 'gist_heat_r',
'gist_ncar', 'gist_ncar_r', 'gist_rainbow',
'gist_rainbow_r', 'gist_stern', 'gist_stern_r',
'gist_yarg', 'gist_yarg_r', 'gnuplot', 'gnuplot2',
'gnuplot2_r', 'gnuplot_r', 'gray', 'gray_r', 'hot',
'hot_r', 'hsv', 'hsv_r', 'inferno', 'inferno_r',
'jet', 'jet_r', 'magma', 'magma_r', 'nipy_spectral',
'nipy_spectral_r', 'ocean', 'ocean_r', 'pink',
'pink_r', 'plasma', 'plasma_r', 'prism', 'prism_r',
'rainbow', 'rainbow_r', 'seismic', 'seismic_r',
'spring', 'spring_r', 'summer', 'summer_r', 'tab10',
'tab10_r', 'tab20', 'tab20_r', 'tab20b', 'tab20b_r',
'tab20c', 'tab20c_r', 'terrain', 'terrain_r', 'turbo',
'turbo_r', 'twilight', 'twilight_r',
'twilight_shifted', 'twilight_shifted_r', 'viridis',
'viridis_r', 'winter', 'winter_r'

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

In a spreadsheet, enter the formula " $=\text{LN}(\text{current price}/\text{original price})$ ".

Changes over Time4

- As an alternative to percentage returns, log returns can be used.
- In some scenarios, they are easier to handle and therefore often preferred in a financial context. ** see next slide
- Figure 8-3 (next slide) shows the cumulative log returns for the single financial time series.
- This type of plot leads to some form of *normalization*:

```
In [19]: rets = np.log(data / data.shift(1)) ❶
```

```
In [20]: rets.head().round(3) ❷
```

```
Out[20]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUR=
Date									
2010-01-01	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2010-01-04	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.006
2010-01-05	0.002	0.000	0.000	0.006	0.018	0.003	0.003	-0.035	-0.003
2010-01-06	-0.002	-0.002	-0.003	-0.018	-0.011	0.001	0.001	-0.010	0.003
2010-01-07	-0.002	-0.010	-0.010	-0.017	0.019	0.004	0.004	-0.005	-0.007

	XAU=	GDX	GLD
Date			
2010-01-01	NaN	NaN	NaN
2010-01-04	0.021	NaN	NaN
2010-01-05	-0.001	0.010	-0.001
2010-01-06	-0.002	0.018	0.016
2010-01-07	-0.006	-0.005	-0.006

```
In [21]: rets.cumsum().apply(np.exp).plot(figsize=(10, 6)); ❸
```

❶

Calculates the log returns in vectorized fashion.

❷

A subset of the results.

❸

Plots the cumulative log returns over time; first the `cumsum()` method is called, then `np.exp()` is applied to the results.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

* Log Returns vs Simple Returns1

Simple Return Formula: $r_i = \frac{P_t}{P_{t-1}} - 1$

Log Return Formula: $r_i = \ln(P_t/P_{t-1})$

$$\ln\left(\frac{P_t}{P_{t-1}}\right) \approx \frac{P_t}{P_{t-1}} - 1 = \frac{P_t - P_{t-1}}{P_{t-1}}$$

Log returns can be easily converted back into simple returns. To get simple returns out from the log returns, you can easily do it by applying the exponential function.

Assignment Project Exam Help

<https://tutorcs.com>

Simple returns are asset-additive: Portfolio return is the weighted average of the stocks in the portfolio.

WeChat: cstutorcs

$$R_{p,t} = \sum_{i=1}^n w_{i,t} R_{i,t}$$

****** One of the advantages is additivity over time, which does not hold true for simple percentage changes/returns. This will be covered in lecture 10 in further detail.

Log returns are not asset-additive. The weighted average of log returns of individual stocks is not equal to the portfolio return. In fact, log returns are not a linear function of asset weights. In comparison, if simple returns are used then the portfolio return is the weighted average of assets in that portfolio. So, one of the advantages of simple return is that it can be used where portfolios are formed and portfolio returns have to be calculated because of its asset-additive property.

Log returns are time-additive: The logarithmic return of an asset over a period of t to T is the sum of all logarithmic returns between the t and T. In other words, the log return over n periods is merely the difference in log between initial and final periods. This is an advantage because the sum of a normally distributed variable is also normally-distributed.

* Log Returns vs Simple Returns2

- See **additional** optional spreadsheet provided **Log>Returns-Part1 and 2.xlsx**. This will be covered in more detail in Week 10. It is provided to purely demonstrate the difference between calculating simple and log returns.

- You will notice that that we take the log of percentage change.

- But take log?

- The reason for this is that log of the returns is time additive.

- That is,
- If r_{13} is the returns for time between t_3 and t_1 .
- r_{12} is the returns between t_1 and t_2 and
- r_{23} is the returns between t_2 and t_3 .
- Then, $\log(r_{13}) = \log(r_{12}) + \log(r_{23})$

For example:,

If $p_1 = 100$, $p_2 = 110$ and $p_3 = 120$, where p_1 is price of stock in time 1

Then:

$$\log(r_{12}) = \ln(p_2/p_1) = \ln(110/100) = 9.53\%$$

$$\log(r_{23}) = \ln(120/110) = 8.7\% \text{ and}$$

$$\log(r_{13}) = \log(r_{12}) + \log(r_{23}) = 9.53 + 8.7 = 18.23\%, \text{ which is same as } \ln(120/100).$$

This means a log change of +0.1 today and then -0.1 tomorrow will give you the same value of stock as yesterday. This is not true if you simply compute percentage change.

It is **common** practice in portfolio optimization to take log of returns for calculations of covariance and correlation

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Cumulative Total Return

- The cumulative return is the total change in the investment price over a set time
- Total Cumulative Return =
$$\frac{(\text{Current Price of Security}) - (\text{Original Price of Security})}{\text{Original Price of Security}}$$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs
- Total Cumulative Return from Simple Returns = (PRODUCT(1+All Simple Returns))-1
- Total Cumulative Return from Log Returns = (EXP(SUM(All log returns))) – 1
- These calculations will be the same answer. It is a double check.

Changes over Time⁵




Figure 8-3. Cumulative log returns over time

Resampling1

- Resampling is an important operation on financial time series data.
- Usually this takes the form of *downsampling*, meaning that, for example, a tick data series is resampled to one-minute intervals or a time series with daily observations is resampled to one with weekly or monthly observations (as shown in Figure 8-4):

```
In [22]: data.resample('lw', label='right').last().head() ❶
Out[22]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX
Date								
2010-01-03	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2010-01-10	30.282827	30.66	20.83	133.52	174.31	114.57	1144.98	18.13
2010-01-17	29.418542	30.86	20.80	127.14	165.21	113.64	1136.03	17.91
2010-01-24	28.249972	28.96	19.91	121.43	154.12	109.21	1091.76	27.31
2010-01-31	27.437544	28.18	19.40	125.41	148.72	107.39	1073.87	24.62

```
EUR= XAU= GDX GLD
Date
2010-01-03 1.4323 1096.35 NaN NaN
2010-01-10 1.4412 1136.10 49.84 111.37
2010-01-17 1.4382 1129.90 47.42 110.86
2010-01-24 1.4371 1092.50 48.78 107.17
2010-01-31 1.3822 1081.05 49.86 105.96
```

```
In [23]: data.resample('lm', label='right').last().head() ❷
Out[23]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX
Date							
2010-01-31	27.437544	28.1800	19.40	125.41	148.72	107.3900	1073.87
2010-02-28	29.231399	28.6700	20.53	118.40	156.35	110.7400	1104.49
2010-03-31	33.571395	29.2875	22.29	135.77	170.63	117.0000	1169.43
2010-04-30	30.126584	30.5350	22.84	137.10	145.20	118.8125	1186.69
2010-05-31	36.697106	25.8000	21.42	125.46	144.26	109.3690	1089.41

```
.VIX EUR= XAU= GDX GLD
Date
2010-01-31 24.62 1.3862 1081.05 40.72 105.960
2010-02-28 19.50 1.3625 1116.10 43.89 109.430
2010-03-31 17.59 1.3510 1112.80 44.41 108.950
2010-04-30 22.05 1.3295 1178.25 50.51 115.360
2010-05-31 32.07 1.2305 1215.71 49.86 118.881
```

```
[24]: rets.cumsum().apply(np.exp).resample('lm', label='right').last(
      ).plot(figsize=(10, 6)); ❸
```


<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.resample.html>
<https://www.w3resource.com/pandas/series/series-last.php>

❶ EOD data gets resampled to *weekly* time intervals ...

❷ ... and *monthly* time intervals.

❸ This plots the cumulative log returns over time: first, the `cumsum()` method is called, then `np.exp()` is applied to the results; finally, the resampling takes place.

Resampling2



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Figure 8-4. Resampled cumulative log returns over time (monthly)

AVOIDING FORESIGHT BIAS

When resampling, pandas takes by default in many cases the left label (or index value) of the interval. To be financially consistent, make sure to use the right label (index value) and in general the last available data point in the interval. Otherwise, a foresight bias might sneak into the financial analysis.³

***3.** *Foresight bias* — or, in its strongest form, *perfect foresight* — means that at some point in the financial analysis, data is used that only becomes available at a **later** point.

We do not want to use data that does not exist in the future in a backtest.

- The result might be “too good” results, for example, when backtesting a trading strategy.

Rolling Statistics

- It is financial tradition to work with *rolling statistics*, often also called *financial indicators* or *financial studies*. Such rolling statistics are basic tools for financial chartists and technical traders, for example.
- This section works with a single financial time series only:

<https://tutorcs.com>
WeChat: cstutorcs

```
In [25]: sym = 'AAPL'
In [26]: data = pd.DataFrame(data[sym]).dropna()
In [27]: data.tail()
Out[27]:
```

Date	AAPL.O
2018-06-25	182.17
2018-06-26	184.43
2018-06-27	184.16
2018-06-28	185.50
2018-06-29	185.11

An Overview1

- It is straightforward to derive standard rolling statistics with pandas:

- 1 Defines the window; i.e., the number of index values to include.
- 2 Calculates the rolling minimum value.
- 3 Calculates the rolling mean value.
- 4 Calculates the rolling standard deviation.
- 5 Calculates the rolling median value.
- 6 Calculates the rolling maximum value.
- 7 Calculates the exponentially weighted moving average, with decay in terms of a half life of 0.5.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

```
21: window = 10 1
29: data['min'] = data[sym].rolling(window=window).min() 2
30: data['mean'] = data[sym].rolling(window=window).mean() 3
31: data['std'] = data[sym].rolling(window=window).std() 4
32: data['median'] = data[sym].rolling(window=window).median() 5
33: data['max'] = data[sym].rolling(window=window).max() 6
34: data['ewma'] = data[sym].ewm(halflife=0.5, min_periods=window).mean() 7
```

A **half-life** is the time taken for something to halve its quantity.

halflife : *float, str, timedelta, optional*

Specify decay in terms of half-life

$\alpha = 1 - \exp(-\ln(2)/halflife)$, for *halflife* > 0.

If **times** is specified, the time unit (str or timedelta) over which an observation decays to half its value. Only applicable to **mean()**, and halflife value will not apply to the other functions.

- <https://www.investopedia.com/terms/e/ema.asp>
- <https://en.wikipedia.org/wiki/Half-life>
- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.ewm.html>

An Overview2

- To derive more specialized financial indicators, additional packages are generally needed (see, for instance, the financial plots with Cufflinks in “Interactive 2D Plotting”).
- Custom ones can also easily be applied via the `apply()` method.
- The following code shows a subset of the results and visualizes a selection of the calculated rolling statistics (see Figure 8-5):

```
In [35]: data.dropna().head()
Out[35]:
```

	AAPL.O	min	mean	std	median	max
Date						
2010-02-01	27.818544	27.437544	29.580892	0.933650	29.821542	30.719969
2010-02-02	27.979972	27.437544	29.451249	0.968048	29.711113	30.719969
2010-02-03	28.461400	27.437544	29.343035	0.950665	29.685970	30.719969
2010-02-04	27.435687	27.435687	29.207892	1.021129	29.547113	30.719969
2010-02-05	27.922829	27.435687	29.099892	1.037811	29.419256	30.719969

```
ewma
Date
2010-02-01 27.805432
2010-02-02 27.936337
2010-02-03 27.728931
2010-02-04 27.659299
2010-02-05 27.856947
```

```
In [36]: ax = data[['min', 'mean', 'max']].iloc[-200:].plot(
        figsize=(10, 6), style=['g--', 'r--', 'g--'], lw=0.8) ❶
data[sym].iloc[-200:].plot(ax=ax, lw=2.0); ❷
```

❶

Plots three rolling statistics for the final 200 data rows.

❷

Adds the original time series data to the plot.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

An Overview3



Figure 8-5. Rolling statistics for minimum, mean, maximum values

Technical Analysis Example1

- Rolling statistics are a major tool in the so-called *technical analysis* of stocks, as compared to the fundamental analysis which focuses, for instance, on financial reports and the strategic positions of the company whose stock is being analyzed.
- A decades-old trading strategy based on technical analysis is using two *simple moving averages* (SMAs).
- **The idea is that the trader should go long on a stock (or financial instrument in general) when the shorter-term SMA is above the longer-term SMA and should go short when the opposite holds true.**
- The concepts can be made precise with `pandas` and the capabilities of the `DataFrame` object.
- Rolling statistics are generally only calculated when there is enough data given the `window` parameter specification.
- As **Figure 8-6** shows, the SMA time series only start at the day for which there is enough data given the specific parameterization:

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Technical Analysis Example2

- The Simple Moving Average (SMA) is calculated by adding the price of an instrument over a number of time periods and then dividing the sum by the number of time periods.
- The SMA is basically the average price of the given time period, with equal weighting given to the price of each period.

Assignment Project Exam Help

Formula

$$\text{SMA} = (\text{Sum}(\text{Price}, n)) / n$$

Where: n = Time Period

<https://tutorcs.com>

WeChat: cstutorcs

```
In [37]: data['SMA1'] = data[sym].rolling(window=42).mean() ❶
```

```
In [38]: data['SMA2'] = data[sym].rolling(window=252).mean() ❷
```

```
In [39]: data[[sym, 'SMA1', 'SMA2']].tail()
Out[39]:
```

Date	AAPL.O	SMA1	SMA2
2018-06-25	182.17	185.606190	168.265556
2018-06-26	184.43	186.087381	168.418770
2018-06-27	184.16	186.607381	168.579206
2018-06-28	185.50	187.089286	168.736627
2018-06-29	185.11	187.470476	168.901032

```
In [40]: data[[sym, 'SMA1', 'SMA2']].plot(figsize=(10, 6)); ❸
```

❶ Calculates the values for the shorter-term SMA.

❷ Calculates the values for the longer-term SMA.

❸ Visualizes the stock price data plus the two SMA time series.

Technical Analysis Example3



Figure 8-6. Apple stock price and two simple moving averages

Technical Analysis Example4

- In this context, the SMAs are only a means to an end. They are used to derive positions to implement a trading strategy.
- Figure 8-7 visualizes a long position by a value of 1, and a short position by a value of -1.
- The change in the position is triggered (visually) by a crossover of the two lines representing the SMA time series:

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

①

Only complete data rows are kept.

②

If the shorter-term SMA value is greater than the longer-term one ...

③

... go long on the stock (put a 1).

④

Otherwise, go short on the stock (put a -1).

```
In [41]: data.dropna(inplace=True) ①
```

```
In [42]: data['positions'] = np.where(data['SMA1'] > data['SMA2'], ②  
                                         1, ③  
                                         -1) ④
```

```
In [43]: ax = data[['sym', 'SMA1', 'SMA2', 'positions']].plot(figsize=(10, 6),  
                                                             secondary_y='positions')  
          ax.get_legend().set_bbox_to_anchor((0.25, 0.85));
```

Technical Analysis Example5



Figure 8-7. Apple stock price, two simple moving averages and positions

- The trading strategy implicitly derived here only leads to a few trades per se:
- only when the position value changes (i.e., a crossover happens) does a trade take place.
- Including opening and closing trades, this would add up to just six trades in total.

Correlation Analysis

- As a further illustration of how to work with `pandas` and financial time series data, consider the case of the S&P 500 stock index and the VIX volatility index.
- It is a stylized fact that when the S&P 500 rises, the VIX falls in general, and vice versa.
- This is about *correlation* and not *causation*.
- This section shows how to come up with some supporting statistical evidence for the stylized fact that the S&P 500 and the VIX are (highly) negatively correlated.**

WeChat: cstutorcs

** One reason behind this is that when the stock index comes down — during a crisis, for instance - trading volume goes up, and therewith also the volatility. When the stock index is on the rise, investors generally are calm and do not see much incentive to engage in heavy trading. In particular, long-only investors then try to ride the trend even further.

The Data1

- The data set now consists of two financial times series, both visualized in Figure 8-8:

```
In [44]: raw = pd.read_csv('..\\source\\trading\\eikon\\eod_data.csv',  
                          index_col=0, parse_dates=True)

In [45]: data = raw[['SPX', 'VIX']].dropna()

In [46]: data.tail()
Out[46]:
```

	.SPX	.VIX
Date		
2018-06-25	2723.06	15.92
2018-06-26	2723.06	15.92
2018-06-27	2699.63	17.91
2018-06-28	2716.31	16.85
2018-06-29	2718.37	16.09

```
In [47]: data.plot(subplots=True, figsize=(10, 6));
```

Note: The filename in code has been set to a directory path.

①

Reads the EOD data (originally from the Thomson Reuters Eikon Data API) from a CSV file.

The Data2



Figure 8-8. S&P 500 and VIX time series data (different subplots)

The Data3

```
In [48]: data.loc[:'2012-12-31'].plot(secondary_y='.VIX', figsize=(10, 6));
```

①

`.loc[:DATE]` selects the data until the given value DATE.

- When plotting (parts of) the two time series in a single plot and with adjusted scalings, the stylized fact of negative correlation between the two indices becomes evident through simple visual inspection (Figure 8-9):

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Figure 8-9. S&P 500 and VIX time series data (same plot)

Logarithmic Returns1

- As pointed out earlier, statistical analysis in general relies on returns instead of absolute changes or even absolute values.
- Therefore, we'll calculate log returns first before any further analysis takes place.
- Figure 8-10 shows the high variability of the log returns over time.
- For both indices so-called “volatility clusters” can be spotted.
- In general, periods of high volatility in the stock index are accompanied by the same phenomena in the volatility index:

```
In [49]: rets = np.log(data / data.shift(1))
```

```
In [50]: rets.head()
```

```
Out[50]:
```

	.SPX	.VIX
Date		
2010-01-04	NaN	NaN
2010-01-05	0.003111	-0.035038
2010-01-06	0.000545	-0.009868
2010-01-07	0.003993	-0.005233
2010-01-08	0.002878	-0.050024

```
In [51]: rets.dropna(inplace=True)
```



Figure 8-10. Log returns of the S&P 500 and VIX over time

Logarithmic Returns2

- In such a context, the pandas `scatter_matrix()` plotting function comes in
- handy for visualizations. It plots the log returns of the two series against
- each other, and one can add either a histogram or a kernel density estimator
- (KDE) on the diagonal (see Figure 8-11):

```
In [53]: pd.plotting.scatter_matrix(rets, ①  
                                         alpha=0.2, ②  
                                         diagonal='hist', ③  
                                         hist_kwds={'bins': 35}, ④  
                                         figsize=(10, 6));
```

①

The data set to be plotted.

②

The alpha parameter for the opacity of the dots.


③

What to place on the diagonal; here: a histogram of the column data.

④

Keywords to be passed to the histogram plotting function.

Logarithmic Returns3



Note: There is a known bug with this scatter matrix function which displays the upper left quadrant axis to a higher level of decimalisation.

Correlation1

- Finally, we consider correlation measures directly.
- Two such measures are considered: a static one taking into account the complete data set and a rolling one showing the correlation for a fixed window over time.
- **Figure 8- 13** illustrates that the correlation indeed varies over time but that it is always, given the parameterization, negative.
- This provides strong support for the stylized fact that the S&P 500 and the VIX indices are (strongly) negatively correlated:

```
In [56]: rets.corr() ❶
Out[56]:      .SPX      .VIX
      .SPX  1.000000 -0.804382
      .VIX -0.804382  1.000000
```

```
In [57]: ax = rets['.SPX'].rolling(window=252).corr( ❷
      rets['.VIX']).plot(figsize=(10, 6))
      ax.axhline(rets.corr().iloc[0, 1], c='r'); ❸
```

❶ The correlation matrix for the whole DataFrame.

❷ This plots the rolling correlation over time ...

❸ ... and adds the static value to the plot as horizontal line.

Correlation2



Figure 8-13. Correlation between S&P 500 and VIX (static and rolling)

High Frequency Data1

Note: The filename in code has been set to a directory path.

- This chapter is about financial time series analysis with pandas.
- Tick data sets are a special case of financial time series.
- A tick is a measure of the minimum upward or downward movement in the price of a security.
- Frankly, they can be handled more or less in the same ways as, for instance, the EOD data set used throughout this chapter so far.
- Importing such data sets also is quite fast in general with pandas.
- The data set used comprises 17,352 data rows (see also **Figure 8-14**):

```
In [59]: %%time
# data from FXCM Forex Capital Markets Ltd.
tick = pd.read_csv('../source/fxcm_eur_usd_tick_data.csv',
                    index_col=0, parse_dates=True)

CPU times: user 1.07 s, sys: 149 ms, total: 1.22 s
Wall time: 1.1 s

In [60]: tick.info()
<class 'pandas.core.frame.DataFrame'>
Int64Index: 461357 entries, 2018-06-29 00:00:00.082000 to 2018-06-29
20:59:00.607000
Data columns (total 2 columns):
Bid      461357 non-null float64
Ask      461357 non-null float64
dtypes: float64(2)
memory usage: 10.6 MB

In [61]: tick['Mid'] = tick.mean(axis=1) ❶

In [62]: tick['Mid'].plot(figsize=(10, 6));
```

❶

Calculates the `Mid` price for every data row.

High Frequency Data2



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Figure 8-14. Tick data for EUR/USD exchange rate

High Frequency Data3

- Working with tick data is generally a scenario where resampling of financial time series data is needed.
- The code that follows resamples the tick data to five-minute bar data (see **Figure 8-15**) (next slide), which can then be used, for example, to backtest algorithmic trading strategies or to implement a technical analysis:

```
In [63]: tick_resam = tick.resample(rule='5min', label='right').last()

In [64]: tick_resam.head()
Out[64]:
```

	Bid	Ask	Mid
2018-06-29 00:05:00	1.156500	1.156501	1.156500
2018-06-29 00:10:00	1.15671	1.15672	1.156715
2018-06-29 00:15:00	1.15725	1.15727	1.157260
2018-06-29 00:20:00	1.15720	1.15722	1.157210
2018-06-29 00:25:00	1.15711	1.15712	1.157115

```
In [65]: tick_resam['Mid'].plot(figsize=(10, 6));
```

High Frequency Data4



Figure 8-15. Five-minute bar data for EUR/USD exchange rate

Conclusion

- This chapter deals with financial time series, probably the most important data type in the financial field.
- `pandas` is a powerful package to deal with such data sets, allowing not only for efficient data analyses but also easy visualizations, for instance.
- `pandas` is also helpful in reading such data sets from different sources as well as in exporting the data sets to different technical file formats.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Further Resources

Good references in book form for the topics covered in this chapter are:

- McKinney, Wes (2017). *Python for Data Analysis*. Sebastopol, CA: O'Reilly.
- VanderPlas, Jake (2016). *Python Data Science Handbook*. Sebastopol, CA: O'Reilly.

<https://tutorcs.com>

WeChat: cstutorcs